

PromptABI: Static Verification for the Discrete Interface Layer of LLM Systems

PromptABI Contributors

June 5, 2026

Abstract

Large language model applications increasingly fail at a layer that is neither neural nor traditionally tested: the discrete interface layer made of chat templates, tokenizer metadata, special tokens, stop policies, structured output schemas, tool definitions, provider request envelopes, and framework truncation rules. PromptABI is a static and differential verifier for that layer. Its core claim is deliberately narrow and therefore practical: before loading model weights or running inference, a developer can prove that important prompt and protocol states are possible, impossible, ambiguous, or unsafe under the exact artifacts that will be used in production.

This paper presents the repository as it now stands and the research system it is designed to become. The implementation already contains a typed Python package, a deterministic command line workflow, a public embedding API, a typed artifact model, a stable diagnostic contract with text, JSON, and SARIF renderers locked by structured snapshots, upstream issue/fix/workaround links, position-aware source spans for configs and JSON artifacts, offline version-pinned artifact loading, a tokenizer abstraction and differential harness validated against real Hugging Face `tokenizers`, `tiktoken`, and `SentencePiece` libraries, deterministic finite automata, finite-state transducers with composition and projection, executable specifications that replay witnesses, bounded product-language laws, and SMT proof objects, a declarative static contract language with a canonical `.pabi` parser, formatter, migration tool, linter, JSON lowering path, CLI check mode, and reviewable examples for LLM apps, fine-tuning datasets, RAG systems, and evaluation harnesses, now compiled into source-mapped finite obligations whose diagnostics preserve human rule names, rule-line spans, solver models, and proof or counterexample classification, lowered to a finite-contract solver with compressed lazy product witnesses, incremental DFA minimization, indexed transducer composition/projection, solver timeouts, constraint-sliced finite enumeration, reusable automata/transducer/SMT counterexample shrinking for minimal strings, token paths, and finite assignments, and a deterministic solver-query cache keyed by normalized Z3 formulas, artifact hashes, supported-fragment metadata, and solver-version fingerprints, deterministic solver-replay JSON files that rerun reduced SMT-backed findings without loading original artifacts, datasets, prompts, or provider credentials, byte-level tokenizer memoization for repeated normalization, a dependency-aware verification-session scheduler with deterministic parallel check execution and a per-session analysis cache, a typed plugin registry for artifact loaders, checks, provider adapters, grammar backends, template dialects, truncation policies, solver encodings, and diagnostic renderers, CLI-loadable module plugins, entry-point plugin discovery, and mode-aware plugin checks that preserve deterministic scheduling, an incremental monorepo verifier that maps changed local artifact paths to check dependencies, reuses cached unchanged-check diagnostics, recomputes downstream dependents, and fails closed by recomputing cache misses rather than silently dropping findings, and a dependency graph reporter that renders configured artifacts, missing artifact-kind surfaces, checker prerequisites, and shared analysis resources as text, JSON, or GitHub-compatible Mermaid diagrams, a generated public-API stability manifest with deprecation metadata, Markdown/JSON API docs, and downstream plugin compatibility tests, focused tests, examples, a fixture corpus layout, a curated ten-family seed corpus, a deterministic corpus manifest pipeline, a Hugging Face `tokenizer_config.json` chat-template parser, a tokenizer/config drift checker that compares real tokenizer revisions for special-token ID changes, added-token changes, normalizer changes, chat-template hashes, BOS/EOS flags and IDs, and generation stop-policy deltas, deterministic PromptABI lockfiles that pin verified artifact hashes, upstream revisions, library versions, provider fixture versions, supported-fragment metadata, and expected diagnostic fingerprints for CI enforcement, signed verification bundles that package config hashes, lockfile state, diagnostics, witness hashes, bounded artifact excerpts, solver metadata, and reproducibility hashes under a local HMAC signature for audit trails, plus supply-chain diagnostics for missing hashes, license metadata, trusted-source annotations, and non-reproducible remote revisions, a dedicated `promptabi verify-training` workflow for training manifests, dataset declarations, tokenizer and chat-template pins, packing windows, loss masks, redaction policy, trained-on-invalid-interface checks for impossible roles, malformed tool calls, invalid JSON targets, and unreachable stop examples, synthetic-data generator preflight checks that reject finite generator plans able to emit forbidden roles, invalid schema outputs, malformed tool-call envelopes, or truncation cases that drop required regions before any generated row is admitted, streaming-shard verification that samples JSONL chunks without materializing private corpora, proof-carrying dataset shard sidecars that bind local verification summaries to streamed shard SHA-256 hashes and hash-only counterexample fingerprints, offline data-loader adapter fixtures for Hugging Face Datasets, JSONL, Parquet metadata, Axolotl, LLaMA-Factory, TRL, OpenRLHF, and internal formats, training-to-serving bridge facts, and corpus-metadata drift checks that compare training baselines against model-card facts, tokenizer/template pins, and serving lockfile summaries while keeping witnesses structural or hash-only rather than raw private dataset material, a first-class evaluation-harness verifier that loads benchmark prompt manifests and proves, over finite declared facts, that provider/model metadata, eval-tokenizer drift pins, prompt templates, prompt variables, few-shot roles/budgets, multi-turn history truncation, system-prompt survival, role alternation, tool availability, JSON answer parsers, schemas, stops, and private answer-key, grader-rubric, and chain-of-thought fields agree with the model/provider contract before a benchmark score is trusted, including bounded cross-provider compatibility reports that compare OpenAI-compatible, Anthropic-style, Gemini-style, Bedrock, local vLLM, llama.cpp, and Ollama eval adapters for request/response fields, tool argument encoding, parallel tool support, stop policies, and structured-output modes, a prompt-pack artifact model, `prompt-pack-contracts` check, package-manager-style `promptabi prompt-pack lock` workflow,

and **promptabi prompt-pack upgrade** gate for reusable prompt libraries, plus a public **promptabi prompt-pack registry** manifest that publishes non-sensitive proof hashes, diagnostic fingerprints, supported-fragment counts, and reproducible artifact metadata without raw template, tool, or stop contents, plus signed **promptabi prompt-pack provenance** manifests that bind reviewed registry metadata, contract hashes, proof hashes, package checksums, versions, and diagnostic baselines under a local HMAC key so downstream teams can verify template, schema, and tool-definition bundles without seeing raw prompt-pack contents, and a **promptabi prompt-pack mirror** workflow that copies local private pack artifacts into checksum-verified, network-free enterprise mirrors and revalidates them offline before use, validating, pinning, and proving upgrade preservation for exported templates, role-boundary and budget hashes, tool schemas, stop policies, supported model/provider families, package hashes, and diagnostic baselines against downstream application artifacts, plus compositional prompt-pack diagnostics that reuse library-level guarantees when an app adds bounded RAG chunks, extra context, tools, and framework truncation policies, proving required pack regions remain declared, must-survive, and not truncated under the same token-budget model used by RAG checks, plus an agent-framework adapter that turns LangChain/LlamaIndex-style runtime prompt assembly specs into ordinary PromptABI configs and proves the selected reusable prompt pack still matches dynamic segments, tools, stops, and provider family choices, a curated **promptabi gallery** that reruns real pinned configurations, emits compatibility badges, summarizes actual solver backends rather than merely advertised solver capability, and explains abstentions plus accepted risks, a **promptabi doctor** preflight that inspects Python and PromptABI version metadata, optional tokenizer/grammar/Z3 dependencies, a live Z3 smoke query, cache writability, first-party and requested plugin registries, supported backend capabilities, config validity, artifact path health, and aggregated setup problems before CI or local verification, lightweight proof-sketch certificates for role-boundary non-forgeability, stop overreachability, tokenizer/grammar emptiness, must-survive budgets, and finite SMT contracts, with executable checks that replay reported counterexamples through the same bounded abstraction and real tokenizer adapters when available, while downgrading to sketches or abstentions when no replayable object exists, deterministic diagnostic clustering that groups related findings by rule, declared root cause, artifact edge, provider behavior, or shared minimized witness digest while preserving unclustered-finding coverage counts for safe triage, deterministic fix-suggestion ranking that orders candidate repairs by safety, downstream compatibility, blast radius, and whether they alter user-visible prompt behavior, a **promptabi fix** engine that can preview or apply low-risk lockfile refreshes, explicit special-token-map declarations, unsupported-fragment annotation files, and reviewer documentation stubs while re-running verification after writes, and can separately produce guarded, non-mutating high-risk previews for template, schema, stop-policy, parser, or truncation changes with before/after witness obligations, review guardrails, and explicit prompt-behavior risk labels, rich diagnostic witness records that carry the human trace plus machine-readable rendered strings, exact token IDs, role-region offsets, parser states, solver assignments, truncation decisions, and ranked minimal fixes, so a CI failure can be replayed by tools, explored in a web-free static HTML report with role-region overlays, token-boundary views, and solver-assignment tables, and audited without reverse-engineering a plain-text explanation, plus opt-in witness privacy modes for text, JSON, SARIF, GitHub annotations, and HTML that leave diagnostic fingerprints, token IDs, source spans, and structural offsets stable while transforming witness payload strings and witness artifact paths into position-preserving redactions, hash-only evidence, or structural length/line/digest summaries, a reusable GitHub Action and **promptabi github-action** runner that restore analysis caches, enforce lockfiles for app configs, run **verify-training** for training-manifest pull requests before fine-tuning jobs start, skip unrelated diffs by watching manifests, JSONL shards, tokenizer/template pins, and configured PromptABI inputs, write markdown job summaries, emit workflow annotations, and upload GitHub-code-scanning SARIF, plus **promptabi pre-commit** local workflows that install PATH-independent git hooks and fail closed on staged schemas, templates, tokenizer configs, tool definitions, training manifests, PromptABI configs, and lockfiles before merge, and a minimal VS Code extension that shells out to the same deterministic CLI to publish inline diagnostics, rerun quick checks, open selected explanations, and preview rendered or tokenized witness traces, an opt-in **-local-summary** usage analytics path that appends only sanitized command names, exit codes, durations, diagnostic counts, and aggregate artifact/check counts to local JSONL while **promptabi usage privacy** explicitly states that prompts, schemas, configs, constraints, witnesses, artifact contents, paths, and network sends are excluded, a **promptabi usage metrics** export that aggregates rule IDs, solver outcomes, artifact categories, and measured check runtimes without raw prompt content, a **promptabi dashboard** team-risk view that rolls up live diagnostics, accepted suppressions, solver abstentions, drift warnings, corpus regressions, and local JSONL trend deltas across configs without exposing prompt contents, a checked security model that documents structural prompt-interface risks, non-goals, secret-handling behavior, artifact privacy, solver-input privacy, provider fixture safety, suppressions, and a responsible disclosure workflow, deep conceptual documentation for the transducer formalism, tokenizer/template composition, Z3/SMT-backed finite contracts, role-boundary non-forgeability, stop reachability, grammar emptiness, and must-survive budget proofs, with tests that tie the documented examples and API names back to executable code, stack integration guides for LangChain RAG, LlamaIndex agents, vLLM OpenAI-compatible servers, llama.cpp/Ollama local servers, Hugging Face Transformers, OpenAI-compatible providers, LiteLLM routers, MCP tools, custom agent frameworks, and training/fine-tuning data pipelines, with guide tests that verify referenced repository artifacts and run representative CLI commands, policy and suppression files that match diagnostics by stable fingerprints, artifacts, or paths; require reviewed justifications, owners, accepted-risk statements, expiration dates, and **witness_digest** proofs that the suppressed counterexample is unchanged; enforce repository severity thresholds; reject stale witness drift; and render audited suppression diagnostics without letting stale or unjustified exceptions weaken CI, a **promptabi minimize** workflow and public API that shrink failing templates, JSON Schemas, stop-string sets, message sets, finite solver-constraint lists, provider fixtures, and formal DFA/FST/SMT counterexamples under failure-preserving or globally minimality-certified predicates into compact upstreamable repros, an upstream bug-report generator that turns a selected real diagnostic into a sanitized markdown issue containing artifact versions and hashes, reproduction commands, expected and actual behavior, linked patches, fixed versions, local workarounds, redacted length-limited witness traces, minimized input summaries when supplied, and privacy guarantees suitable for filing with tokenizer, template, framework, or provider maintainers, a maintainer refresh workflow that regenerates model-artifact manifests, structured-schema manifests, provider fixture pack manifests, real-bug benchmark manifests, portable expected-diagnostic baselines, corpus diffs, and release notes from real verification output, a contributor-infrastructure gate that validates structured GitHub issue templates, a label manifest with good-first-issue and area labels, MkDocs-linked plugin/checker/corpus author guides, and CI path triggers through **promptabi**

contribute validate, a **promptabi init** workflow that writes verifiable starter contracts and local fixture stubs for OpenAI tools, Hugging Face local models, vLLM OpenAI-compatible servers, llama.cpp, LangChain RAG, LlamaIndex agents, and custom JSON Schema outputs, a **promptabi diff** workflow that compares two full model/provider/framework configurations and emits contract-breaking tokenizer, provider, framework, artifact-set, check-surface, and context-budget changes with structured witnesses and migration guidance, plus a **promptabi version-gate** workflow that classifies each real diff diagnostic as patch-safe, minor-breaking, or major-breaking under deployment policy overrides, a **promptabi launch-assets** workflow that regenerates comparison tables against prompt linters, schema validators, constrained decoders, tokenizer diff tools, and generic static analyzers; a Mermaid architecture diagram; a terminal demo script with a valid GIF placeholder; SVG benchmark charts; bug-gallery markdown; Hacker News/GitHub positioning copy; and a manifest whose headline counts are sourced from the live benchmark, evaluation, beta, and real-bug replay code paths rather than hand-edited launch claims, a stop-policy parser that normalizes OpenAI-compatible request snapshots, Hugging Face generation configs, llama.cpp/Ollama options, vLLM sampling params, LiteLLM passthrough params, and common framework wrapper kwargs into one typed artifact model with source spans, a stop-tokenizer analyzer that records byte/string/token alignments, multi-token stops, invalid stop-token IDs, normalization-sensitive stops, prefix/suffix collisions, and special or added token interactions against real tokenizer adapters, a fixture-backed stop differential harness that replays OpenAI-compatible, Hugging Face, llama.cpp-style, and vLLM-compatible stop traces through small CPU-only simulators, and a bounded stop-overreachability checker that constructs serialized witnesses for supported schema/tool string fields, nested JSON, markdown code fences, XML-like tool calls, and provider tool-call envelopes, including exact line/column firing points, parser states at truncation, and the malformed or prematurely accepted structures delivered to downstream parsers, a bounded symbolic executor for the supported Jinja chat-template fragment, a bounded role-boundary non-forgeability checker that detects unsanitized content and role fields capable of rendering control delimiters, special tokens, assistant prefixes, or tool-call sentinels, while recognizing explicit JSON, escape, and delimiter-safe wrapper filters, minimized role-forgery witnesses that include the malicious input field, rendered prompt excerpt, byte-level tokenization, and the exact forged boundary, a role-boundary model that decomposes symbolic renders into system, user, assistant, tool, developer, function, and provider-specific structural regions, a concrete renderer differentially checked against Hugging Face **apply_chat_template** behavior across the seed corpus, a delimiter- collision regression suite covering ChatML, Llama header tokens, Mistral instruction tags, XML-ish tool tags, markdown role fences, and fine-tune hash headers, a real-world bug corpus that reduces public llama.cpp, vLLM, and Hugging Face Transformers reports into synthetic offline fixtures for Phi-style role delimiter forgery, ChatML special-token injection, Qwen/Gemma-style tool argument truncation, streaming stop interruption, array-of-object tool-argument leakage, multi-function XML boundary confusion, and XML tool-call parser-boundary stops, a replayable real-bug benchmark manifest that extends those reductions with tokenizer-differential drift, provider migration, structured-output parser mismatch, RAG truncation, and supervised target-role-alignment failures under one deterministic CPU-only command while preserving upstream links, fixed versions, and local workarounds, a labeled evaluation runner that expands real-bug cases and a maintained eval-fixture pack for stop-string, parser, truncation, role-boundary, and tokenizer-mismatch bugs into scoped precision/recall, abstention rate, runtime, peak-memory, witness-quality, solver-quality, and real-library differential-agreement metrics, an offline beta-program replay that runs open-source-style case-study configs, records upstream issue evidence, checks tuned-absent false-positive rules, and gates abstention messages on actionable witnesses and setup suggestions, a deterministic generated-property suite that exercises every public checker entry point over real safe, unsafe, ambiguous, unsupported, satisfiable, unsatisfiable, and abstaining artifacts without adding runtime dependencies, a grammar ingestion layer that normalizes JSON Schema, regex, EBNF, Outlines, xgrammar, llguidance, and PromptABI grammars into typed rules, terminals, source spans, supported-fragment metadata, and explicit abstentions, plus a JSON Schema path that normalizes objects, arrays, enums, consts, required fields, numeric and string constraints, type and composition unions, bounded local references, recursion-limit abstentions, and additional-property semantics, then compiles the supported subset into grammar IR, parser states, DFA witnesses, source maps, and real **jsonschema** validator round-trips, replays a CPU-only grammar differential corpus with hand-labeled Outlines, xgrammar, llguidance, lm-format-enforcer, Guidance, Instructor, and Pydantic-derived semantics through the CLI as agreement, mismatch, or abstention diagnostics, a structured-output parser compatibility checker that replays bounded grammar/schema witnesses and declared runtime samples against JSON, JSON Schema, XML-ish tool-call, markdown-fence, and custom-delimiter parser models while separating parser-broader from grammar-broader disagreements and explicitly marking the result as heuristic evidence rather than language equivalence, a labeled structured-schema corpus with open-source-agent reductions, anonymized production-style contracts, synthetic stress cases, tool-definition bundles, per-entry PromptABI configs, deterministic hashes, and manifest generation, a tool/function schema ingestion layer that normalizes OpenAI tools, Anthropic tools, LangChain tools, Pydantic model schemas, TypeScript-style functions, MCP tools, and provider tool-call envelopes into source-mapped typed function schemas, plus a bounded tool-call serialization checker that compares recorded provider, template, stop-policy, and parser contracts for tool names, JSON-object versus JSON-string arguments, escaping guarantees, tool-call IDs, parallel calls, streaming chunk assembly, and tool-call delimiters, a bounded provider migration checker that compares recorded OpenAI, Azure OpenAI, Anthropic, Gemini, Bedrock, Together, Groq, Ollama, llama.cpp server, vLLM OpenAI server, and LiteLLM contracts for request/response field loss, tool argument encoding, tool-call IDs, parallel-call support, streaming chunks, stop sequences, context limits, structured-output modes, error envelopes, and LiteLLM routing, and now takes the bounded tokenizer x grammar product to prove whether compiled witnesses survive real tokenizer encode-normalize-decode assumptions, plus a bounded tokenizer x grammar ambiguity check that validates JSON lexical variants and detects tokenizer-path conflicts, decoded-text conflicts, whitespace and Unicode byte aliases, and added-token aliases, a secret-free provider fixture pack corpus that records and replays request and response shapes, tool-call encodings, stop behavior, streaming deltas, error envelopes, context/output limits, and labeled edge cases for OpenAI, Anthropic, Gemini, Bedrock, vLLM OpenAI-compatible servers, and LiteLLM, with deterministic manifests, replay hashes, source-mapped diagnostics, and validation that rejects credential-like fields, a context-window budget model that combines named prompt segments, must-survive flags, declared or tokenizer-derived token counts, output reserves, tool-token reserves, generation-prompt tokens, and model-specific special-token overhead into bounded survival proofs, compact text/JSON/SARIF budget visualizations, and web-free static HTML reports with interactive witness, prompt, chart, diff, and corpus sections against real framework truncation policies, including minimized counterexamples when

required regions can be dropped, an enterprise-readiness surface that checks strict no-network declarations, prompt-pack mirror manifests and mirrored-file digests, private artifact indexes, and internal provider fixtures through the shared secret scanner, plus organization-wide policy packs that layer required checks, severity thresholds, supported proof-mode fragments, solver-timeout caps, local-summary privacy rules, and approved provider-fixture digests into the existing severity/suppression pipeline with restrictive multi-pack merging, severity overrides that run before threshold gating, and declarative solver timeout/memory/network sandbox posture without pretending to enforce an OS sandbox from inside the process, a dedicated RAG chunking compatibility check that verifies tokenizer identity, chunk-boundary stability, citation survival, overlap accounting, metadata inflation, prompt-template overhead, and retrieval payload limits against the same packed prompt model, and a Z3-backed static-contract layer that derives finite SMT obligations from real templates, tokenizers, special-token maps, stop policies, tools, provider configs, truncation configs, prompt segments, and training/fine-tuning manifests, proving or refuting prompt-budget survival, stop/control-token exclusion, role-region non-forgeability, tool-schema required-parameter satisfiability, tool/provider compatibility, supervised target-role alignment, and observed supervised-span containment across rendered role regions, tokenized bounds, preserved dataset-packing boundaries, BOS/EOS delimiters, loss-mask selection, and exclusion of declared user, tool, retrieval, or preference source ranges from supervised targets, chosen/rejected preference-pair agreement over prompt prefixes, role layouts, tokenizer pins, mask policies, response starts, truncation, and preserved packing boundaries, including a CPU-only RLHF/DPO fixture pack with synthetic and anonymized clean, prompt-mismatched, mask/truncation-defective, and role-layout/packing-defective manifests, plus tokenizer, chat-template, model-card, serving-lockfile, role-delimiter, special-token, generation-prompt, tool-format consistency, and finite invalid-interface summaries that reject misspelled canonical roles, malformed trained tool-call envelopes, invalid structured-output targets, and stop examples unreachable under the declared training interface, plus synthetic-generator contracts that catch role, schema, tool-call, and truncation failures before private or large generated datasets are materialized, with concrete models, explicit proof classifications, or minimized unsat cores, a **promptabi solver replay** maintainer path that verifies such reduced solver obligations from portable JSON and treats solver backends, query hashes, and Z3 versions as provenance rather than brittle equality requirements, a minimized **promptabi corpus smt-benchmark** line with satisfiable, unsatisfiable, timeout-prone, and unsupported prompt-interface obligations replayed against the real finite/Z3 solver, a CPU-only benchmark suite, documentation, and contribution guidance, plus **promptabi release compatibility-audit** and **promptabi release drift-bisect** gates that replay pinned corpus/provider/ grammar/parser/framework fixtures and binary-search local tokenizer, template, schema, provider, or framework artifact histories to identify the first contract-regressing revision—abstaining rather than certifying unpinned upstream releases—and a **promptabi release readiness** gate that checks version metadata, changelog, README command anchors, stable CLI paths, GitHub Action and release workflows, MkDocs navigation, seed-corpus manifests, formal/Z3-backed check registration, real-bug replay, the reproducibility bundle, and this paper preprint against live repository code before declaring a 1.0 shipment. The CLI now also includes **promptabi explain**, which replays verification, deduplicates repeated logical diagnostics, selects one finding by fingerprint, rule, or index, and renders a tutorial explanation with source snippets, the checked formal property, proof-mode descriptions, witness steps, likely production symptoms, and concrete fixes in text or JSON. Diagnostic messages now also expose localization-ready catalog keys and strict placeholder metadata, with **promptabi diagnostics catalog** emitting a deterministic English source catalog from real verification output while existing CI renderers continue to display stable English text. The same diagnostic objects now feed **promptabi diagnostics lsp**, a language-server-style `textDocument/publishDiagnostics` protocol that groups config and artifact findings by document URI, converts source spans to zero-based editor ranges, preserves fingerprints, suggestions, witnesses, and proof modes in diagnostic data, and publishes empty diagnostic arrays for watched local schemas, templates, stop policies, training manifests, and artifact bundles so editor shims can clear stale inline findings deterministically. The **promptabi contract** workflow migrates deprecated syntax with automata/solver behavior notes, parses, canonicalizes, and composes human-authored `.pabi` contracts. It emits the lowered JSON artifact, enforces stable rule names, and merges organization policies, prompt-pack contracts, app configs, and training-manifest contracts with explicit precedence: role permissions narrow by intersection, required regions and forbidden delimiters tighten by union, and lower layers cannot silently downgrade policy severity. Rule names and source spans survive into **static-contracts** diagnostics, so a Z3/finite counterexample or proof points back to the governing `.pabi` line. A concrete VS Code extension now consumes that payload directly: it registers quick-check, explanation, and witness-preview commands, maps PromptABI spans to native diagnostics, attaches code actions to findings, and previews the exact render/tokenize witness steps produced by real checks rather than a separate editor-only analysis. The adjacent **promptabi bug-report** command uses the same selector and verification path but emits an upstream-ready markdown issue rather than a local tutorial, including upstream status, fixed versions, and local workarounds, so the strongest evaluation unit—a maintainer-actionable bug report—is generated from the exact diagnostic object exercised in CI. The remaining checkers fit into that surface rather than requiring a rewrite. We describe the transducer-oriented formalism, the artifact model, the diagnostic contract, the deployed benchmark lines, and the planned evaluation against real tokenizer, template, structured-output, provider, and budget failures. The diagnostic layer now records whether a finding is sound, complete, bounded, Z3-backed, heuristic, or an abstention, making each claim explicit about its proof obligations. The result is a path toward a tool that is useful in CI and credible as a systems paper: it catches failures that arise before any model logit is relevant.

1 Introduction

PromptABI is a static verifier for the discrete application binary interface of LLM systems: the files and adapters that translate application state into model visible bytes, tokens, structured-output languages, tool calls, and parser inputs. This interface is now complex enough to fail independently of model quality. A template can let user text forge a role header, a stop sequence can truncate a valid tool argument, a tokenizer revision can reinterpret a sentinel, or a framework budget policy can drop a required system segment while ordinary unit tests still pass.

The contribution is to make those failures reproducible, local, and explicit. PromptABI models supported fragments of templates, tokenizers, grammars, stop policies, provider fixtures, and budget policies as finite automata, finite-state transducers, differential replays, and bounded SMT obligations. It reports concrete witnesses when a bad state is reachable, proof metadata when a finite contract excludes one, and abstentions when an artifact escapes the modeled fragment. The paper therefore argues for a deliberately narrow systems claim: many high-impact LLM application bugs are structural interface bugs, and they can be found before inference.

2 Motivation

Modern LLM applications are built from small, fragile protocol objects. A chat template decides which bytes separate roles. A tokenizer decides whether those bytes are ordinary content or special tokens. A stop policy decides where output is truncated. A tool schema decides which JSON fragment is valid. A provider adapter decides whether a tool call appears as a message, a delta, a function object, or a model-specific sentinel. A framework budget policy decides which messages, retrieved chunks, and tool definitions survive when the prompt exceeds the context window.

These objects compose into a pipeline:

```
messages -> chat template -> byte/string prompt -> tokenizer -> token stream
          -> constrained decoder / stop logic / tool parser -> application parser
```

The failure modes are concrete. A user field can contain a string that renders as an assistant delimiter. A stop string can fire inside a valid JSON string, making the application parse a malformed prefix. A tokenizer update can change special-token IDs while tests still pass because they never inspect the token stream. A RAG prompt can preserve the latest user message while silently dropping the safety instruction or the tool definition that gives the message meaning. These are not model-behavior questions. They are contract questions.

PromptABI treats those contracts as first-class verification targets. It does not promise semantic prompt-injection resistance. It does not ask whether a model will follow an instruction. It asks whether the interface representation admits a structural state that the developer did not intend. That boundary makes the tool useful without GPUs, private model access, or nondeterministic generation.

3 Current Repository Artifact

The current repository milestone establishes the skeleton needed for serious verification work. It uses a Python `src/` layout, declares package metadata in `pyproject.toml`, exposes the `promptabi` console entrypoint, ships `py.typed`, and includes focused tests for the public API, artifact model, diagnostic contract, configuration loader, and CLI behavior. The implementation is intentionally small, but it is not a paper mockup: it loads real JSON configuration, discovers configs from subdirectories, accepts explicit artifact overrides, normalizes artifact paths, builds typed artifact bundles, records source spans for config declarations and parsed JSON artifact fields, materializes local and metadata-only remote artifacts, validates sha256 pins, audits licenses, trusted-source annotations, and reproducible remote revisions, adapts real tokenizer backends into stable encode/decode metadata, extracts Hugging Face chat-template structure from `tokenizer_config.json`, derives sanitizer-aware role-boundary regions over bounded rendered prompt patterns, emits minimized role-forgery witnesses with attacker input, rendered excerpt, token evidence centered on the forged marker, exact boundary offsets, and explicit sanitizer metadata, extracts automata, transducer, and solver witnesses for finite formal contracts and independently checks those witnesses with executable specifications, including role-boundary counterexamples re-substituted into bounded renders and re-tokenized against real adapter behavior, tracks check dependencies over artifact kinds, reuses cached tokenizer adapters, tokenizer/grammar analysis reports, budget reports, and finite static-contract query results keyed by normalized formulas plus artifact/fragment/solver fingerprints, emits replay files containing only reduced finite obligations, expected logical results, artifact hashes, and precise privacy metadata, parallelizes independent embedded checks, serializes shared tokenizer and solver resources, renders the same scheduler metadata as an artifact-to-check dependency graph for configured and missing tokenizer, template, schema, tool, stop-policy, truncation-policy, provider, prompt-segment, and training-manifest surfaces, exposes an incremental monorepo plan that resolves explicit or `git diff -relative` changed paths to local artifact kinds, forces full runs when config files change, reuses JSON-hydrated diagnostics for unchanged checks, recomputes cache misses to preserve CI correctness, and records structured reuse/cache-miss diagnostics in the same result stream, fails closed on malformed JSON or UTF-8 artifacts, corrupt fixture archives, partial provider snapshots, invalid schemas, missing or incompatible tokenizer backends, and unsupported solver fragments by emitting source-mapped loader diagnostics or explicit abstentions, emits deterministic typed diagnostics in text, JSON, SARIF, GitHub annotations, and static HTML with structured snapshot coverage, exposes an embedding API for custom checks, typed plugin registries, plugin artifact-loader hooks, plugin renderers, and typed results, ingests JSON Schema, regex, EBNF, Outlines, xgrammar, llguidance, and PromptABI grammar artifacts into typed rule and terminal summaries with source spans and abstention issues, normalizes supported JSON Schema objects, arrays, unions, references, constraints, and additional-property policies into a typed tree, compiles that tree into bounded grammar IR with parser states and DFA witnesses, and round-trips generated witnesses through the real `jsonschema` validator, checks grammar-backend differential fixtures against hand-authored accepted and rejected strings for Outlines, xgrammar, llguidance, lm-format-enforcer, Guidance, Instructor-style, and Pydantic-derived schemas, abstaining rather than overclaiming for fragments outside local replay semantics, compares structured-output grammar/schema artifacts with declared application parser models for JSON, JSON Schema, XML-ish tool calls, markdown fences, and custom delimiters using bounded witnesses and runtime samples, distinguishing parser-broader acceptance that can admit invalid application states from grammar-broader acceptance that can generate runtime-parse failures, checks tokenizer x grammar ambiguity by pairing schema-valid JSON lexical variants with real tokenizer encode/decode behavior to find conflicting token paths, decoded-text collapses, Unicode or whitespace byte aliases, and added-token aliases, checks recorded tool-call serialization contracts for provider/schema/parser agreement over names, argument encodings, escaping, IDs, parallel calls, streaming fragments, chat-template tool rendering, and stop delimiters, checks recorded provider migration contracts across all current first-party provider families for request, response, tool, stop, streaming, context, structured-output, error, and routing compatibility, validates provider fixture packs that capture secret-free request/response shapes, tool-call encodings, stop behavior, streaming deltas, error envelopes, and edge-case limits for offline replay, and emits cross-provider evaluation reports over real provider fixtures so benchmark claims can distinguish comparable adapter surfaces from missing request fields, stop-policy drift, unsupported parallel tools, or absent structured-output modes, ships paired end-to-end buggy/fixed mini applications for tool calling, JSON structured output, RAG truncation, provider migration, supervised training-data alignment, trained-on-invalid-interface examples, and transform-attributed training source leakage, with each pair verified by the exact checker that catches the bug and by the fixed contract that removes it, models context-window budgets from named prompt segments, must-survive flags, declared or real-tokenizer token counts, output reservations, reserved tool tokens, generation-prompt tokens, model-specific special-token overhead, bounded truncation-policy survival proofs with minimized failing segment sets, and machine-readable budget visualizations that expose spans, dropped fields, and survival guarantees, ships a dependency-free notebook repr API whose rich HTML and plain-text views wrap the same tokenizer encoders, concrete and symbolic chat-template renderers, stop-policy reachability analyzer, lazy automata product search, finite SMT/contract solver, and token-budget visualization objects used by CLI checks, with escaped self-contained markup and explicit proof-mode labels for teaching, debugging, and paper artifact exploration, derives Z3-backed finite static contracts over cross-artifact obligations such as required-prompt-token budget survival, stop strings colliding with special or added control tokens, controlled prompt regions containing role/control boundary markers, tool schemas requiring parameters absent from their declared properties, tool-provider family mismatches, and training target roles outside the serving chat-template role universe, plus supervised training spans that no longer align with the observed rendered role region, token bounds, packed example boundary, or loss mask, source-contribution leakage into target spans, chosen/rejected preference-pair divergence before the compared response, and tokenizer/template drift between fine-tuning preparation, training, evaluation, and serving pins, and lets maintainers replay those obligations from fixture JSON while validating stored SAT witnesses directly against the deserialized formula, statically checks evaluation harnesses so prompts, few-shots, grading parsers, application parsers, tokenizer snapshots, and stops use the provider/template/schema

contract verified for serving, loads reusable prompt packs as first-class artifacts, locks them with package-manager-style contract hashes, and checks that their exported templates, expected roles, required tools, stop sequences, and supported model/provider families remain compatible with the app that imports them, then activates a bounded composition layer when the app supplies retrieval chunks or a framework truncation policy: required prompt-pack regions must be declared as must-survive prompt segments, the existing budget proof must not drop them, retrieval payloads must carry runtime token bounds, and any added non-retrieval roles are surfaced as explicit guarantee-boundary warnings rather than silently absorbed into the pack proof, while the public registry format converts the same verified state into contract hashes, hash-only template/tool/stop proofs, diagnostic fingerprints, supported-fragment counts, and reproducible source metadata suitable for publishing outside a private prompt library, and the local mirror workflow turns reviewed private packs into content-addressed files plus a manifest that detects checksum, size, and path-escape drift without network access, returns stable policy-controlled process exit codes, writes and enforces deterministic verification lockfiles that fail closed on artifact, config, provider-fixture, library-version, supported-fragment, or diagnostic-baseline drift, and adds an opt-in artifact-provenance check for hash, license, trusted-source, and reproducible-download supply-chain gates, adds enterprise-ready offline controls: local mirror manifests can be pinned by sha256, private indexes must be local and source-allowlisted, internal provider fixtures are parsed with the same redaction scanner used by first-party fixture packs, organization policy packs require named checks, cap solver timeouts, restrict supported proof-mode fragments, forbid local usage summaries when the privacy policy demands it, approve internal provider fixtures by digest, and merge restrictively across multiple packs, rule-level severity overrides are applied before CI gating, remote artifact locations are rejected under strict no-network mode, and solver sandbox settings are reported honestly as declarative timeout, memory, and network posture rather than as unproven runtime isolation, ships cross-version CI over supported Python releases, operating systems, optional tokenizer/grammar/Z3 dependency slices, and pinned backend versions, plus a repository-local GitHub Action that composes cache restoration, lockfile enforcement, training-data PR gates, changed-artifact detection, SARIF upload, workflow annotations, and markdown summaries around the same verifier, and installs a local pre-commit workflow that gates staged PromptABI-relevant schemas, templates, tokenizers, tools, training manifests, configs, and lockfiles while failing closed when staged content differs from the working tree, records optional telemetry-free local command summaries and local metrics exports with no artifact paths, prompt text, schema bodies, solver constraints, witnesses, or network transport, documents the local trust boundary, structural security claims, non-goals, solver privacy, provider-fixture secret rejection, sanitized upstream reporting, and responsible disclosure path in a MkDocs-linked security model, adds MkDocs concept chapters for formal automata/transducers, tokenizer-template composition, Z3/SMT finite contracts, role-boundary non-forgeability, stop reachability, grammar emptiness, must-survive budgets, and proof-sketch certificates, and verifies those chapters against live symbols and executable examples in the test suite, adds a MkDocs-linked integration guide that maps PromptABI artifacts and checks to LangChain, LlamaIndex, vLLM, llama.cpp/Ollama, Hugging Face Transformers, OpenAI-compatible servers, LiteLLM, MCP tools, internal agent frameworks, and training pipelines, and tests that the documented stack scaffolds, fixture paths, and smoke commands execute against the real CLI, packages releases as checked PyPI distributions, standalone executables, Docker images, signed/attested GitHub assets, changelogs semantic-version tags, and GHCR container images, compares baseline and current PromptABI configs with `promptabi diff`, emitting text, JSON, or SARIF diagnostics for removed checks, missing artifacts, tokenizer/config drift, provider-contract regressions, framework truncation regressions, and top-level context-budget reductions, scaffolds common stacks with `promptabi init` and verifies every generated config through the same loader, provider replay, tool-schema, parser, budget, RAG, stop, and role-boundary checks used elsewhere in the test suite, expands a selected diagnostic with `promptabi explain` into a local debugging narrative that includes artifact snippets, the formal property, witness trace, likely production symptom, and fix suggestions while handling repeated logical findings deterministically, publishes `promptabi matrix` as a machine-readable compatibility table that aggregates every check’s emitted rule IDs into explicit sound/complete/bounded/Z3/heuristic/abstaining modes across tokenizer families, template dialects, grammar backends, providers, frameworks, and the currently modeled training-manifest surface, applies repository policy files after real diagnostics are produced so accepted risks are owned, expiring, fingerprint-stable, and bound to unchanged witness digests rather than blanket ignores, runs CPU-only performance benchmarks over tokenizer, template, grammar, stop, SMT, budget, corpus, and cache paths, mutates chat templates, tokenizers, stop policies, schemas, grammars, tools, truncation configs, and SMT encodings to discover newly introduced contract violations, and writes a paper reproducibility package with frozen fixture hashes, stable expected tables, solver-version pins, environment metadata, and one-command regeneration scripts, plus a maintainer artifact bundle with model/corpus/provider manifests, expected diagnostic baselines, corpus diffs, generated release notes, and a launch asset bundle whose comparison table, diagram, benchmark chart, public bug gallery, demo script, and concise positioning are checked by focused tests, adds a `promptabi corpus bug-gallery` path that converts replayed real-bug benchmarks into sanitized artifacts, minimized structural witnesses, root-cause explanations, fixes, and upstream patch links, and validates contribution infrastructure—issue templates, labels, contributor guides, docs navigation, and CI triggers—through the same local CLI/API path used by maintainers.

The repository layout is designed to scale:

Path	Purpose
src/promptabi	Typed package, artifact model, formal automata, finite-state transducers, finite-contract core, layered contract composition, offline loaders, CLI, configuration loading, diagnostic model, renderers, notebook reprs, scheduler, analysis cache, and verification session.
src/promptabi/incremental.py	Monorepo incremental verification planner, git changed-path collection, per-check diagnostic cache, dependency-closure selection, cache reuse diagnostics, and fail-closed cache-miss recomputation for promptabi verify -changed-path/-changed-from-git .
src/promptabi/dependency_graph.py	Deterministic artifact/check graph builder and text, JSON, and Mermaid renderers for configured artifacts, missing artifact-kind requirements, checker prerequisites, and shared analysis resources exposed through promptabi graph .
src/promptabi/plugins.py	Typed extension registry for artifact loaders, dependency-aware checks, diagnostic renderers, and capability metadata covering provider adapters, grammar backends, template dialects, truncation policies, and solver encodings.
src/promptabi/prompt_packs.py	Prompt-pack compatibility and composition checker for reusable prompt libraries, including downstream role/tool/stop/provider matching, required-region survival under truncation, bounded RAG payload checks, package lock/upgrade gates, public registry manifests with non-sensitive proof hashes, signed provenance manifests, and checksum-verified local mirror manifests for air-gapped private prompt-pack indexes.
src/promptabi/agent_frameworks.py	Dynamic agent-framework adapter that loads prompt-pack assembly specs, validates selected templates against runtime segments and dynamic context sources, renders previews, and writes ordinary PromptABI configs verified by the real prompt-pack checker.
src/promptabi/api_stability.py	Generated public-API manifest, stable/provisional/deprecated symbol taxonomy, deprecation decorator, Markdown/JSON API reference renderers, and baseline compatibility comparison for downstream plugin authors.
src/promptabi/contributor_validation.py	Dependency-free validation for issue templates, label manifests, plugin author docs, checker design rules, corpus contribution rules, MkDocs navigation, and CI contributor-test gates exposed as promptabi contribute validate .
editors/vscode	Minimal VS Code extension shim that invokes promptabi diagnostics lsp and promptabi explain , publishes inline diagnostics, contributes quick-check/code-action commands, and renders witness previews from real verification payloads.
tests	Focused tests for package API and CLI determinism.
examples/minimal	A typed artifact config, messages file, tool schema, and answer schema used by tests and docs.
examples/prompt-packs	A reusable support prompt-pack fixture and downstream app contract proving prompt-library imports preserve role, tool, stop, model-family, bounded RAG, and truncation-survival assumptions.
examples/agent-frameworks	Safe and intentionally drifted dynamic agent prompt-pack specs plus a runnable generator that materializes runtime prompt assembly into promptabi.json before deployment.
examples/role-boundary	Paired unsafe and sanitized ChatML-style tokenizer_config.json artifacts that exercise the real role-boundary non-forgeability CLI check.
examples/stop-policies	A vLLM sampling-parameter fixture that exercises real stop-policy artifact loading and normalization.
examples/end-to-end	Paired buggy and fixed mini apps for tool-calling serialization, JSON parsers, citation-preserving RAG, provider migration, training role alignment, and a GPU-free chat SFT quickstart.
examples/gallery	Curated real-world-style configurations that rerun pinned OpenAI-tool, ChatML role-boundary, training/static-contract, RAG-budget, and accepted-risk policy examples with deterministic proof, badge, and risk summaries.
fixtures and paper_artifact	Curated seed corpus, real-world bug reductions, grammar differential cases, labeled structured-schema/tool-calling fixtures, secret-free provider packs, and a generated paper artifact containing fixture hashes, expected tables, solver pins, and regeneration commands.
benchmarks and promptabi fuzz	CPU-only benchmark entrypoints plus deterministic mutation fuzzing for tokenizer analysis, symbolic template execution, grammar products, stop checks, SMT contracts, budget modeling, corpus verification, cache reuse, and cross-artifact violation discovery.
promptabi launch-assets	Evidence-backed launch material: comparison table, architecture diagram, demo GIF source script, SVG benchmark chart, public bug gallery, positioning copy, and a manifest generated from real evaluation, benchmark, beta, and real-bug replay code paths.

This milestone matters because later checkers must not accrete as one-off scripts. They need stable input objects, stable diagnostics, stable renderers, stable examples, and stable tests from the beginning. PromptABI now has that surface and a model that names each interface object explicitly.

4 Core Abstraction

The unifying abstraction is that LLM interface systems are compositions of finite-state transducers over byte, string, token, and structured alphabets, with partial parsers at the boundaries. A chat template maps structured messages to a string. A tokenizer maps strings to token streams and back, subject to special tokens, normalization, byte fallback, and added-token behavior. A constrained decoder or grammar describes a language of valid outputs. A stop policy cuts a prefix from that language. An application parser maps the resulting bytes back to an abstract syntax tree or rejects them.

Many high-value questions become language questions:

Property	Question
Reachability	Can a declared delimiter, stop sequence, or control token be produced under the selected tokenizer and grammar?
Over-reachability	Can the same sequence appear inside user-controlled content, a JSON string, a tool argument, or a markdown block?
Emptiness	Is the product of tokenizer assumptions and grammar constraints empty, so no valid output can be generated?
Non-forgability	Can an attacker-controlled field render as role or provider structure after template expansion?
Must-survive	Does a required prompt segment remain after the framework truncation policy is applied?

The abstraction is not used to pretend arbitrary Python, arbitrary Jinja, and arbitrary provider behavior are finite-state. Instead, PromptABI classifies each check using a small contract taxonomy: *sound* means the diagnostic is not reported unless the modeled violation exists; *complete* means every violation inside the supported fragment is found; *bounded* means exactness is restricted to declared finite limits; *Z3-backed SMT* means the finite symbolic contract is lowered to Z3 when the optional solver is present; *heuristic* means the result is useful evidence rather than a proof; and *abstaining* means the checker deliberately declines unsupported cases. That taxonomy is crucial: it lets the tool be useful where it is precise and honest where it is not. The same taxonomy is now executable metadata: the compatibility matrix derives check rows from scheduler dependencies and per-rule diagnostic modes, so, for example, stop-tokenizer analysis is shown as a mixture of sound unreachable-stop results, heuristic collision evidence, and abstentions, while static contracts are shown as bounded Z3-backed SMT over all artifact kinds, including finite supervised target-role alignment and rendered/tokenized/loss-masked span containment for training manifests. The central proof obligations also have an executable specification layer: automata witnesses are replayed transition-by-transition, product automata are checked by bounded language equivalence rather than by trusting product state names, transducer witnesses are replayed with epsilon-aware input/output labels, SAT contract results are re-evaluated against their finite domains and named constraints, UNSAT cores are checked for deletion-minimal inconsistency on small domains, and UNKNOWN results are required to remain explicit abstentions rather than masquerading as proofs.

Theorem 1 (Bounded counterexample soundness). *For any check classified as sound and bounded, if PromptABI emits a concrete witness for a modeled violation, replaying the witness through the same loaded artifacts and bounds reaches the reported unsafe state.*

Theorem 2 (Honest supported-fragment boundary). *For template, grammar, provider, solver, and parser behavior outside the declared supported fragment, PromptABI must emit an abstention or a heuristic diagnostic rather than a proof-labeled safety or unsafety result.*

5 Artifact Model

PromptABI uses a typed artifact model because the same file can participate in multiple checks. A tokenizer config is needed for token-boundary analysis, special-token drift, chat-template parsing, and prompt-budget accounting. A tool definition is needed for role-boundary analysis, structured-output compatibility, provider migration checks, and budget survival. A diagnostic must therefore identify what artifact it refers to without coupling the renderer to a particular loader.

The public API now provides `ArtifactKind`, `ArtifactLocation`, `ArtifactProvenance`, `ArtifactBundle`, and typed artifacts for tokenizers, chat templates, special-token maps, stop policies, schemas, grammars, tool definitions, prompt segments, provider configs, framework truncation configs, reusable prompt packs, and training/fine-tuning manifests that explicitly model supervised and preference datasets, system-message policies, role-label canonicalization, loss-mask construction, packing windows, and pinned chat-template versions. Preference-pair contracts now include branch-level prompt hashes, response spans, tokenizer pins, mask policies, truncation flags, and packed-example IDs, which is enough to replay RLHF/DPO data-preparation failures without materializing private examples. Training manifests can also carry offline data-loader adapter fixture summaries: PromptABI reads local Hugging Face Datasets metadata, JSONL samples, Parquet metadata stubs, Axolotl configs, LLaMA-Factory dataset-info files, TRL and OpenRLHF preference rows, and internal redacted schemas, then checks declared fields, splits, counts, and non-sensitive fingerprints before treating the loader surface as covered. It also provides `ArtifactRef`, `SourceSpan`, `WitnessTrace`, `Diagnostic`, `VerificationConfig`, `VerificationSession`, and `VerificationResult`. The API surface is now governed by a generated `PublicApiManifest`: stable plugin and embedding symbols are marked separately from provisional research surfaces, deprecated callables carry machine-readable replacement/removal metadata and emit `DeprecationWarning`, and `promptabi api-docs` renders the same manifest as Markdown or JSON for downstream compatibility baselines. These types establish the compatibility contract:

- artifacts have stable kinds, names, local paths or URIs, provenance, optional versions, revisions, hashes, licenses, sources, and trusted-source metadata;
- kind-specific fields capture stop strings, schema dialects, grammar types, tool names, prompt segment requirements, provider families, truncation strategies, and training dataset/mask/packing/template contracts;
- spans are one-based, validated at construction time, and source-mapped from PromptABI configs plus tokenizer, schema, tool, provider, and fixture JSON artifacts;
- witnesses are explicit traces, not prose hidden in the message string;
- diagnostics have rule IDs, severities, artifacts, spans, witnesses, suggestions, ranked fix metadata, and explicit verification modes;
- results serialize to deterministic dictionaries for JSON and snapshots.

The config loader exercises the model by accepting both legacy path strings and typed artifact objects, resolving local paths relative to the config file, attaching a declaration span to each artifact, and retaining URI-backed artifacts without subjecting them to local existence checks. This small behavior prevents a future class of flaky tests and confusing CLI output: diagnostics refer to the same canonical artifact no matter where the command is launched from, while missing-artifact and pinning findings point at the exact config entry that caused the run to inspect that artifact.

6 Formal Verification Core

PromptABI now includes the first executable formal core rather than only a roadmap for one. The core has two deliberately finite pieces. The first is a deterministic finite automata library over explicit alphabets. It constructs literal languages, finite-language tries, and prefix-closed literals; supports completion, complement, union, intersection, and difference; checks emptiness; extracts shortest witnesses with state-based breadth-first search; and serializes automata products deterministically. This is enough to express early reachability facts such as whether a stop string intersects a grammar prefix, whether a user-controlled field can equal a role delimiter, and whether removing forbidden control strings leaves any accepted user language.

The second piece is a finite contract solver. A contract declares variables over Boolean, enum, integer-range, and bounded-string domains, then composes named constraints with equality, disequality, integer comparisons, conjunction, disjunction, negation, implication, membership, string length, and substring containment. When **z3-solver** is available the adapter lowers those finite domains and constraints to Z3, including explicit domain restrictions for bounded strings and per-query timeout policies. When Z3 is absent, the same contract runs through deterministic finite-domain enumeration with constraint-variable slicing, partial assignment pruning, explicit assignment limits, and abstention rather than false proofs when a bound is reached. Both backends return a stable result object with **sat**, **unsat**, or **unknown** status, backend provenance, counterexample assignments for satisfiable bad states, explicit counterexample/proof/abstention conclusions, checked-assignment counts, and a separate solver-budget classification: proved within budget, bounded by an assignment cap, timed out, approximated by an explicitly declared abstraction, or abstained because the formula exceeded the supported encoding. Z3 cores are deletion-minimized against tracked named constraints while finite-domain restrictions remain background assumptions, so the proof object shown to a developer is not merely the first raw solver core.

Solver calls now have a first-class deterministic cache surface. The cache key normalizes the finite problem, sorted named constraints, Z3 s-expressions when available, explicit assignment and timeout budgets, artifact hashes, supported fragment metadata, and solver-version fingerprints. It can be used in memory or round-tripped as JSON, and a cache miss recomputes the query rather than silently reusing stale proof objects. This matters for privacy-preserving CI: teams can reuse expensive static-contract conclusions without keying solely on prompt names or persisting private prompt text.

The formal core is now accompanied by **promptabi.specs**, a machine-checkable reference layer used by the tests and paper artifact. Its DFA spec replays state paths for shortest reachability witnesses and checks eager intersection, union, and difference products by enumerating bounded words over the union alphabet and comparing against the mathematical language operation. Its FST spec replays transition labels, including epsilon insertions and deletions, and reconstructs the accepted input/output pair from the label path. Its finite-contract spec validates SAT assignments by domain membership plus constraint re-evaluation, validates UNSAT cores by independent finite enumeration on toy-sized domains, and validates the sound abstention boundary by ensuring UNKNOWN results carry no counterexample or proof object. These specifications are deliberately slower and simpler than the optimized algorithms; their purpose is to make the paper’s central claims executable against real objects produced by the repository.

The automata layer now optimizes the exact same properties rather than adding a separate approximate engine. It can find intersection witnesses lazily without materializing full products, compressing tokenizer-sized alphabets by equivalent left/right transition signatures at each product state. Eager union and difference remain on completed products so implicit sink-state semantics are preserved. Product construction minimizes completed operands incrementally before forming reachable products, and witnesses are reconstructed from predecessor maps instead of storing full paths in every queue entry. The transducer layer indexes right-hand transitions by input label during composition and reuses source transition indexes during projection closures, removing the nested scans that would otherwise dominate realistic render/tokenize/decode relations. The byte-level tokenizer also memoizes portable normalization results with a bounded cache, which removes repeated Unicode normalization work in corpus-wide stop, grammar, and budget checks without affecting encode/decode metadata.

The repository now uses that core as a real static-verification layer, not just as a toy solver. The **static-contracts** check derives independent cross-artifact obligations from loaded PromptABI artifacts. Required prompt segments and framework truncation configs become an integer SMT query asking whether required tokens can exceed the finite input budget. Stop policies, special-token maps, and tokenizer added-token metadata become a bounded string/enum query asking whether an application stop is also a model control token. Prompt segments, chat-template role universes, stop policies, special tokens, and tokenizer added-token metadata become a role-region non-forgeability query that asks whether controlled user, tool, function, or retrieval regions can contain a finite boundary marker. Source-mapped tool schema metadata becomes a required-parameter satisfiability query that proves every required argument is declared under the tool’s properties, or returns a missing-parameter model. Tool definitions and provider configs become a provider-family compatibility query. Training manifests and chat-template role declarations become a target-role alignment query for supervised data pipelines, and preference manifests become a DPO/RLHF branch-equivalence query: the chosen and rejected branches must share prompt hashes, role layouts, tokenizer versions, mask policies, response start positions, truncation status, and preserved packing boundaries before the compared response. Satisfiable bad-state queries produce concrete model assignments and are reported as counterexample witnesses. Unsatisfiable queries produce proof diagnostics with the solver backend, minimized core, and a safety or incompatibility classification plus the solver-budget reason. Unsupported, timeout-limited, assignment-capped, approximated, or insufficiently finite inputs are separated explicitly rather than widening into a claim the implementation cannot justify.

The core now adds the missing relational layer: finite-state transducers over explicit input and output alphabets. A transducer transition carries an input label, an output label, or epsilon on either side. This directly models render-like insertions, tokenizer-like recodings, decode-like normalization, and provider-envelope rewrites. The implementation accepts finite relations, literal string mappings, identity relations, exact composition over a shared middle alphabet, input and output projection back

to deterministic automata, shortest input/output witnesses with path labels, and a conservative over-approximation that pairs the independent input and output projections when only a relation envelope is justified. Composition preserves the exact/over-approximate provenance flag so later diagnostics can distinguish proof from bounded approximation.

This split is important for the research claims. Automata should prove language and reachability properties; transducers should model interface mappings and their compositions; SMT should prove finite symbolic compatibility over discrete pipeline facts such as roles, escaping flags, bounded content, token budgets, schema preconditions, and stop-policy states. The fallback enumerator is not a replacement for richer solving, but it makes the supported fragment executable, testable, and privacy-preserving even in minimal CI environments.

The current tests exercise the core against real code rather than sketches. They prove that automata products extract a shortest stop-string witness, that difference excludes a role delimiter while preserving safe user strings, that transducer composition reconstructs a render-then-tokenize witness, that projections recover accepted input and output languages through epsilon labels, that projection-based over-approximation intentionally admits independent pairs, that a finite contract emits a concrete role-forgery counterexample, that contradictory role constraints produce a minimal unsat core, that Boolean compositions are recognized as unsatisfiable, that the installed Z3 backend returns the same kind of concrete assignment for a satisfiable finite contract, that raw Z3 cores are minimized deterministically before becoming diagnostics, and that CLI static-contract findings classify proof-of-safety and concrete counterexample outcomes in the witness trace. The executable-spec tests add an orthogonal benchline: they reject malformed DFA paths, detect deliberately wrong product-language claims, replay composed render/tokenize transducer witnesses, re-check concrete finite-contract assignments, confirm minimal unsat cores, and prove that artificial solver cutoffs surface as bounded or timed-out unknowns rather than false safety. This is the first formal benchline on which role-boundary, stop-reachability, grammar-emptiness, and must-survive budget checkers can be built.

6.1 Proof-Sketch Certificates

The repository now packages the proof obligations above as first-class certificates rather than leaving them implicit in diagnostics. The `promptabi.proof_sketches` module defines a small certificate object with a property identifier, theorem statement, supported-fragment assumptions, lemmas, evidence, outcome, and executable checks. The CLI exposes the registered theorems through `promptabi.proofs`, while the Python API exposes concrete certificate builders for the five core families.

The design is intentionally conservative. A role-boundary certificate validates that a real `RoleBoundaryNonforgeabilityReport` is supported and, for each counterexample, checks that the malicious field is exactly the marker, that the marker appears in the rendered excerpt and tokenizer evidence, and that stored offset lengths match the marker without confusing global rendered offsets for excerpt-relative offsets. A stop-overreachability certificate checks the exact prefix split: the configured stop occurs at the firing offset, the delivered prefix equals the output before that offset, the valid prefix includes the stop, and truncation removes a required suffix. A grammar/tokenizer certificate is marked as a sketch unless the caller also supplies the compiled grammar DFA; with that automaton present, the certificate replays the stored grammar witness through `check_dfa_witness` and rechecks that decoded tokenizer output is accepted by the grammar. A must-survive budget certificate checks that proven results drop no required segment, violated results partition required segments into survived and dropped sets with a minimized counterexample, and unknown results carry a reason instead of silently treating unknown token counts as zero. A static-contract certificate delegates SAT assignment and UNSAT core replay to `check_contract_result`, so Z3-backed answers are independently checked against finite-domain semantics.

This layer is not a substitute for mechanized theorem proving, but it closes an important systems gap: every proof-like claim now has a machine-readable theorem, explicit assumptions, an executable validation boundary, and a visible outcome of `proven`, `counterexample`, `abstained`, or `sketch`.

7 Version-Pinned Artifact Loading

Artifact identity is not enough for a verifier. A checker that proves a property about `tokenizer_config.json` must know which bytes, directory manifest, provider snapshot, or upstream revision it inspected. PromptABI therefore adds an offline loader layer that turns typed artifact references into deterministic `LoadedArtifact` summaries. The loader does not contact third-party services and does not require tokenizer, provider, or grammar libraries. It validates the reproducibility envelope first, so deeper semantic checkers inherit stable inputs.

The implemented loader covers the artifact sources that appear earliest in real LLM interface workflows:

Source	Loaded summary
Local files	Streamed sha256 digest, byte size, pin validation, JSON source-map extraction, and mismatch diagnostics.
Hugging Face refs	Offline repository ID, optional artifact path, revision extraction, and warnings for missing or movable revisions.
Tokenizer directories	Stable manifest over sorted relative file names, file sizes, and content hashes; requires tokenizer metadata.
GGUF stubs	Header-level GGUF magic, version, tensor count, metadata-key count, and streamed file hash without loading weights.
Provider snapshots	Deterministic JSON validation for provider name, captured request/response or streaming shape metadata, provider fixture pack sections, and optional stop-trace fixtures for offline differential replay, with parse-error line and column diagnostics.
Tool definitions	Source-mapped ingestion for OpenAI function tools, Anthropic <code>input_schema</code> tools, LangChain <code>args_schema</code> tools, Pydantic model JSON Schemas, TypeScript-style function arrays, MCP <code>inputSchema</code> tools, and provider tool-call response envelopes, recording tool names, argument encodings, closed parameter schemas, enum/constraint paths, IDs, streaming chunks, and ingestion issues.
Stop-policy configs	Deterministic extraction of stop strings, stop token IDs, EOS behavior, source family, ignored empty stops, and source paths from OpenAI-compatible, Hugging Face, llama.cpp/Ollama, vLLM, LiteLLM, and wrapper shapes.
Prompt segments, RAG chunks, and budgets	Source-mapped prompt segment files from explicit <code>segments</code> objects or OpenAI-style message arrays, merge behavior for config-declared must-survive regions, retrieval chunk metadata for document IDs, tokenizer names, source and packed boundaries, citations, overlaps, metadata costs, template wrapper overhead, payload limits, plus framework context-window artifacts with output, tool, generation-prompt, and special-token reservations.
Schema and grammar files	JSON Schema normalization and bounded compilation for objects, arrays, required fields, enums, consts, constraints, unions, bounded local references, recursion-limit abstentions, and additional properties; compiled metadata records grammar rules, parser states, DFA witness size, source spans, and real <code>jsonschema</code> round-trip acceptance, while tokenizer x grammar emptiness diagnostics prove whether bounded witnesses survive tokenizer encode/decode assumptions and tokenizer x grammar ambiguity diagnostics expose JSON lexical variants, Unicode escapes, whitespace, tokenizer normalization, and added-token paths that collapse or alias structured outputs. Ingestion for regex, EBNF, Outlines, xgrammar, llguidance, and PromptABI grammars records rules, terminals, references, source spans, and unsupported-fragment abstentions.
Fixture archives	Zip/tar member lists and archive-byte hashes without extraction or timestamp-dependent manifests.

Pinning is intentionally behavior-safe. A missing local file is an error because the verifier cannot inspect it. A malformed sha256 or mismatched sha256 is an error because the configured claim is contradicted. An otherwise readable but unpinned artifact is a warning by default, allowing existing projects to adopt strict reproducibility incrementally or fail on warnings in CI. In-memory API artifacts are treated as embedded objects rather than external supply-chain inputs.

This capability is more than bookkeeping. It is the first supply-chain and debuggability boundary for the research system: every future role-boundary, stop-reachability, grammar-emptiness, provider-migration, and budget-survival finding can cite the exact artifact reference, actual hash, revision strength, loader witness, and source span that made the claim reproducible.

8 Chat-Template Parsing and Symbolic Execution

PromptABI now parses Hugging Face `tokenizer_config.json` chat templates into a typed summary instead of treating them as opaque strings. The parser extracts the template source and its JSON source span, declared special tokens from standard tokenizer fields, message and tool field references, loops over `messages` and `tools`, role-condition literals, generation-prompt branches, filters, constants, whitespace-control usage, and unsupported Jinja constructs that require abstention by later symbolic checks.

This is intentionally a parser for verification inputs, not a general-purpose Jinja runtime. It recognizes the bounded fragment used by the seed corpus and records unsupported macros, imports, blocks, complex loops, and runtime globals such as `raise_exception`. The offline loader now recognizes chat-template artifacts backed by `tokenizer_config.json`, validates the template, attaches parsed metadata to the loaded artifact, preserves source spans, and keeps malformed or missing `chat_template` fields as deterministic loader errors. Focused tests run the parser across every curated seed-corpus template and targeted fixtures for tools, filters, whitespace controls, generation prompts, special-token declarations, and unsupported constructs. This creates the concrete input contract needed for bounded symbolic execution and role-boundary non-forgeability checks.

The next layer is now executable. PromptABI includes a position-preserving lexer and bounded symbolic executor for the supported Hugging Face Jinja fragment. It models loops over `messages` and `tools`, `if/elif/else` branches, simple `set` bindings, constants, special-token variables, message and tool fields, filters as symbolic wrappers, explicit whitespace controls, and Hugging Face’s `trim_blocks/lstrip_blocks` structural whitespace behavior. The result is not a single rendered prompt but a finite set of symbolic paths with rendered patterns, path conditions, loop iteration counts, and explicit abstentions when unsupported constructs or path budgets are reached. Loader metadata records whether symbolic execution was precise and how many bounded paths were produced for a real artifact.

On top of those paths, PromptABI now builds a structural role-boundary model. For each symbolic rendered pattern it identifies message-indexed regions, records the conditions that prove a concrete role such as `system`, `user`, `assistant`, `tool`, `developer`, or `function`, keeps unconstrained role variables symbolic, preserves residual else-branch exclusions, and creates explicit generated-assistant-prefix regions. Region offsets point into the symbolic pattern, content expressions identify the controlled fields, and loader metadata summarizes role coverage and region counts for real `tokenizer_config.json` artifacts. This is the executable substrate for the non-forgeability checker rather than a prose-only design.

The first checker over that substrate is now implemented. It builds a global catalog of structural markers from declared special tokens, literal control fragments, generated assistant prefixes, role headers, and XML-like tool sentinels, then proves for each bounded symbolic path whether a raw user-controlled content field or dynamic role field can render one of those markers as structure. It explicitly flags direct `message['role']` interpolation as role-header forgeability and reports bounded, sound diagnostics with rendered excerpts. The checker is now sanitizer-aware at the exact expression level: JSON encoders such as `tojson`, HTML/XML escape filters, URL or base64 wrapper filters, and sanitizer-preserving `set` bindings are recorded in the role-boundary model and suppress only the findings for the protected expression. Raw interpolation remains forgeable, and Jinja’s `safe` marker is deliberately not treated as escaping.

The regression suite now exercises the known delimiter-collision families against the executable checker rather than against string literals. ChatML cases verify `<|im_start|>`, `<|im_end|>`, assistant prefixes, and raw dynamic role headers. Llama-style cases cover start/end header IDs and end-of-turn tokens. Mistral-style cases cover `[INST]`, `[/INST]`, and EOS sentinels. XML-ish tool cases cover opening and closing tool-call tags, markdown cases cover role-bearing code fences, and custom fine-tune cases cover `### User: / ### Assistant:` headers. The tests also assert that every expected finding carries a minimized malicious input, rendered excerpt, byte-level token evidence, and monotonic boundary offsets.

PromptABI also now has a concrete evaluator for that same supported fragment. It renders real message and tool inputs through the position-preserving lexer and bounded parser, evaluating role conditionals, generation-prompt branches, `messages` and `tools` loops, special-token variables, simple bindings, `loop.last`, and JSON filters without silently approximating unknown expressions. A reusable differential harness compares concrete renders with oracle strings collected from Hugging Face’s real chat-template renderer. The deployed tests exercise every seed-corpus template with role/system variants, generation prompts, Unicode and sentinel-like content, structural whitespace, special-token insertion, empty tool lists, tool-role messages, and a focused tool loop using `tojson`. This closes the loop between the symbolic abstraction and the renderer developers actually run.

9 Tokenizer Abstraction Layer

Tokenizers are not a peripheral convenience in PromptABI; they are one of the central semantic boundaries. A role delimiter, stop string, JSON escape, or tool sentinel is only meaningful after normalization, added-token recognition, special-token handling, byte fallback, tokenization, and decoding are known. The repository now exposes a backend-neutral tokenizer interface that records token IDs, token surfaces, byte spans, special-token flags, added-token flags, normalization steps, decoded text, and round-trip metadata.

The implemented layer covers four practical families:

Backend	Capability
Byte-level	Deterministic byte alphabet with greedy added-token matching, portable Unicode normalization rules, UTF-8 byte spans, and exact decode reconstruction.
Hugging Face <code>tokenizers</code>	Local <code>tokenizer.json</code> and <code>tokenizer-directory</code> loading, real backend encode/decode, backend normalizer metadata, offsets converted to byte spans, and added-token detection.
<code>tiktoken</code>	Installed encoding adapters such as <code>cl100k_base</code> , special-token handling, token-byte metadata, and round-trip checks against the real library.
SentencePiece	Local <code>.model</code> loading, piece/id extraction, decode comparison, and best-effort byte spans where the piece surface can be matched.

This abstraction is deliberately thinner than a tokenizer reimplementaion. Its job is to expose the observable behavior of real libraries in a stable shape that formal checks can consume. The implemented stop-tokenizer checker asks for byte spans, token IDs, token surfaces, decoded text, and special/added-token flags without knowing whether the source was ByteLevel BPE, SentencePiece, or `tiktoken`. A role-boundary checker can see whether a delimiter is an ordinary byte sequence, an added token, or a special token.

The current tests validate this layer against real code rather than mocks. A trained local Hugging Face ByteLevel BPE tokenizer is compared against `tokenizers.Tokenizer.encode` and `decode`. The `tiktoken` adapter is compared against `cl100k_base`, including special-token and byte-fallback behavior. A temporary SentencePiece model is trained and compared against `SentencePieceProcessor`. The byte-level adapter separately proves normalization, added-token, special-token, UTF-8 span, and round-trip behavior. These tests form the first differential benchline for the eventual tokenizer corpus.

9.1 Tokenizer Differential Harness

PromptABI now has a reusable differential harness rather than one-off adapter assertions. Each case records the input text, whether the backend should inject special tokens, whether decoding should skip them, and an oracle expectation collected from the real library. The harness then compares PromptABI's stable abstraction against that oracle across token IDs, token surfaces, decoded text, normalized text, normalization steps, special-token IDs, added-token IDs, required byte spans, and normalized round-trip behavior. Failures are returned as deterministic mismatch records with stable dictionaries suitable for future snapshot tests and corpus reports.

The harness already caught a realistic edge case during development: Hugging Face stores special tokens in its added-token decoder, but downstream structural checks need to distinguish provider/model control tokens from ordinary added tokens such as tool sentinels. The adapter now treats specialness and added-token status as separate flags, so a role-boundary checker can reason about `<s>` and `</s>` differently from an application-defined `<tool>` token. This is the sort of small tokenizer-library semantic difference that later formal checks must get right before their witnesses are credible.

9.2 Tokenizer and Config Drift

PromptABI now treats tokenizer upgrades as contract changes instead of informal release-note review. A tokenizer artifact can point at a local drift baseline via `metadata.drift_baseline_path`; the checker reads both baseline and current revisions from real `tokenizer_config.json`, `tokenizer.json`, `special_tokens_map.json`, and `generation_config.json` files. It compares special-token names, surfaces, and IDs; added-token IDs and specialness; normalizer signatures; chat-template hashes and lengths; BOS/EOS token text, IDs, and insertion flags; and stop strings or EOS stop-token IDs. Breaking changes emit deterministic CLI diagnostics with baseline/current values and revision witness steps; exact matches emit an informational proof that the baseline still matches the current artifact. The focused tests build synthetic but real tokenizer directories and run the checker through `promptabi verify`, so the feature is exercised against the same loader, config, diagnostic, and renderer code used by users.

10 Diagnostic Contract

PromptABI's diagnostic contract is designed for CI, code scanning, embeddings, and regression snapshots. A useful finding is not merely a line of text. It needs a stable rule ID, a severity, a precise artifact reference with provenance, a source span when available, a witness that reproduces the failure, a deterministic fingerprint, a suggested remediation, and a declared verification mode. If two releases emit different diagnostics for the same artifact bundle, that difference should be intentional, reviewable, and testable.

The repository now renders the same diagnostic object in human text, deterministic JSON, and SARIF 2.1.0. This matters early: once a verifier has many rules, unstable output becomes a tax on CI, pull-request review, and paper reproducibility. PromptABI therefore fixes ordering, key names, optional-field behavior, and code-scanning fingerprints before high-volume checkers are introduced. A future role-forgery finding can reuse the same shape:

```
{
  "rule_id": "role-boundary-nonforgeability",
  "severity": "error",
  "fingerprint": "stable-sha256-prefix",
  "message": "user.content can render an assistant delimiter",
  "artifact": {
    "kind": "chat-template",
    "name": "tokenizer_config",
    "revision": "pinned-upstream-revision",
    "sha256": "artifact-hash"
  },
  "span": {"path": "tokenizer_config.json", "start_line": 12, "start_column": 3},
  "witness": {
    "summary": "malicious user content reaches assistant header",
    "steps": [
      {"action": "render template", "input": "user.content", "output": "..."},
      {"action": "tokenize prompt"},
      {"action": "locate forged boundary"}
    ]
  },
  "suggestions": ["Escape user content before inserting role delimiters."],
  "check_modes": ["bounded", "sound"]
}
```

The first implemented rule, **repository-skeleton**, is intentionally informational. It proves the machinery works without claiming a security or compatibility property that has not yet been implemented. That restraint is important for trust: checkers should only claim properties that the code has actually proved. The rule is therefore marked **heuristic**; loader and artifact-existence findings are marked **sound** and **complete** for the local filesystem and bytes they inspect; future automata and solver checkers can claim **bounded**, **z3-backed-smt**, or **abstaining** as their supported fragments require.

The SARIF renderer maps PromptABI severities to SARIF levels, emits a rule table, attaches artifact locations and source regions when available, stores the PromptABI fingerprint in **partialFingerprints**, and includes the same check-mode metadata as JSON and text output. GitHub-specific options now add stable analysis categories, repository-root **uriBaseId** rewriting for absolute local paths, deterministic invocation metadata, PromptABI-owned location fingerprints, and schema-correct SARIF suppressions when diagnostics carry external suppression metadata. A separate **github-annotations** renderer emits escaped GitHub Actions workflow commands, so pull requests can receive inline annotations even before a full SARIF-upload Action is installed. This makes the first CLI workflow immediately compatible with GitHub code scanning while keeping the diagnostic source of truth in Python types rather than renderer-specific structures.

The source-region layer is executable rather than aspirational. PromptABI now builds a small recursive JSON source map over the exact text that is later decoded, converts offsets to one-based line and column spans, and keeps those spans outside generic metadata so JSON and SARIF remain serializable. Config artifact declarations carry spans into the typed artifact model; local JSON artifacts expose field-level spans for tokenizer configs, schemas, tool definitions, provider snapshots, and archived fixture members; malformed JSON loader errors carry the decoder's line and column. This gives early diagnostics the same property future formal witnesses need: a developer can jump directly to the config entry, **chat_template**, schema field, or tool declaration that made the finding possible.

The explanation renderer consumes the same diagnostic object rather than a parallel ad hoc format. It first collapses duplicate logical diagnostics by their stable structured payload, then lets users select by fingerprint, rule ID, or one-based index. A successful explanation includes the diagnostic message, fingerprint, artifact and source span, a bounded local source excerpt when the artifact is available, the formal property associated with the rule, human descriptions of the proof modes, the witness trace, the likely production symptom, and the attached fix suggestions. This makes a code-scanning finding actionable without weakening CI determinism or sending prompts to a service.

The minimization workflow is the bridge from local finding to upstream issue. It accepts JSON cases for templates, schemas, stop strings, message sets, solver constraints, and provider fixtures, canonicalizes unordered surfaces, and then uses deterministic

chunk and element deletion with memoized predicate calls. A candidate is accepted only when the original failure predicate still holds; the tool rejects non-reproducing originals rather than emitting success-shaped garbage. The CLI exposes safe local oracles such as substring preservation and a diagnostic oracle that rewrites a scratch artifact, reruns PromptABI, and checks that a selected rule still fires. The public API accepts richer in-process predicates, so tests can prove that a minimized stop-string set still triggers the real tokenizer collision analyzer rather than merely preserving a token in a string.

11 CLI and Embedding API

The command line entrypoint is:

```
promptabi verify
promptabi verify --config examples/minimal/promptabi.json
promptabi verify --config examples/role-boundary/unsafe.promptabi.json
promptabi verify --config examples/role-boundary/safe.promptabi.json
promptabi explain --config examples/role-boundary/unsafe.promptabi.json --index 1
promptabi explain --config examples/minimal/promptabi.json --rule-id repository-skeleton --format json
promptabi minimize repro.json --keep-substring "<|im_start|>" --format json
promptabi init --stack openai-tools --output-dir .promptabi-demo
promptabi gallery --format text
promptabi verify --artifact schema=schemas/answer.json --fail-on warning
promptabi verify --config examples/minimal/promptabi.json --format json
promptabi verify --config examples/minimal/promptabi.json --format sarif
promptabi verify --cache-dir .cache/promptabi --verbose
python examples/agent-frameworks/dynamic_support_agent.py examples/agent-frameworks/safe.agent-prompt-pack.json
```

In the current milestone the workflow validates repository wiring, local existence, reproducible pins, loader compatibility, source-span reporting, and check-mode reporting. It discovers `promptabi.json` or `.promptabi.json` by walking upward from the current directory unless `-config` is explicit. It accepts repeated `-artifact NAME=PATH_OR_URI` overrides so CI can verify generated or matrix-selected artifacts without rewriting the repository config. It prepares a cache directory selected by `-cache-dir`, `PROMPTABI_CACHE_DIR`, `XDG_CACHE_HOME`, or a user cache fallback. It provides quiet and verbose text modes while keeping JSON and SARIF stable for machines.

`promptabi init` is deliberately not a static README snippet. Each stack template writes a `promptabi.json` plus minimal local artifacts that the real loader can ingest: tool schemas, provider replay packs, stop policies, Hugging Face `tokenizer_config.json` templates, prompt segments, truncation budgets, or JSON Schemas as appropriate. The focused init regression test generates every supported stack and verifies the resulting config through the public CLI, which makes the quickstart path an exercised benchline rather than a hand-maintained example.

The exit-code policy is configurable. The default fails only on error-level diagnostics. CI can instead fail on warnings, fail on any diagnostic, or never fail the process while still emitting machine-readable results. Malformed invocations and invalid configs return code 2, preserving the distinction between a broken workflow and a real artifact-contract finding.

The same behavior is available through the embedding API, which now covers the full first integration loop: loading a config, applying artifact overrides, loading artifacts through the deterministic offline loader, selecting built-in or custom checks, collecting diagnostics, returning a typed result, and rendering that result in the same formats as the CLI. The extension surface is no longer only a third-party promise: PromptABI now ships a default first-party plugin pack for Hugging Face, OpenAI-compatible APIs, vLLM, llama.cpp, LangChain, LlamaIndex, Outlines, xgrammar, llguidance, Pydantic, MCP, and Z3. The pack is web-free and import-light: it registers inspectable provider, grammar, template, truncation, tool-schema, artifact-loader, diagnostic-renderer, and solver capabilities, exposes deterministic first-party URI reference loading, and adds a coverage check that proves the bundled adapters are available inside ordinary verification sessions.

```
from promptabi import CheckContext, Diagnostic, DiagnosticSeverity, run_verification
```

```
def organization_policy(context: CheckContext):
    schema = context.artifact("schema")
    yield Diagnostic(
        rule_id="organization-schema-present",
        severity=DiagnosticSeverity.INFO,
        message=f"loaded {schema.artifact.name}",
        artifact=schema.artifact.to_ref(),
    )

result = run_verification(
    "promptabi.json",
    checks={"organization-schema-present": organization_policy},
    selected_checks=["organization-schema-present"],
)

if not result.ok:
    raise SystemExit(1)
```

The public helpers `create_session`, `load_artifacts`, `collect_diagnostics`, `run_verification`, and `render_result` deliberately mirror the CLI rather than bypassing it. The same embedding surface now exposes `create_bug_report` and `render_bug_report`, allowing downstream CI systems to attach sanitized upstream issues or maintainer tickets without scraping

text output. It also exposes `build_public_api_manifest`, `compare_public_api_manifests`, and `public_api_reference`, so plugin authors can pin a JSON baseline and fail upgrades only when a stable import path or required signature is removed. The new `promptabi.agent_frameworks` helps load dynamic prompt-pack agent specs, validate that selected pack templates expose the required roles and regions, render a deterministic prompt preview, and write a normal `promptabi.json`; safe and intentionally drifted examples are then checked by the same `prompt-pack-contracts` rule rather than a special-case example-only verifier. The gallery command uses the same session path, but changes the unit of presentation: each curated entry is a real config run with resolved/pinned artifacts, deterministic diagnostic counts, compatibility badges, actual finite solver backend provenance such as `z3` versus `finite-enumeration`, and visible accepted-risk or abstention explanations. The point is not that the current checker is deep. The point is that all later checkers have an integration path. The built-in plugin pack, a pre-commit hook, a GitHub Action, a notebook visualization, and a provider-fixture replay harness can all depend on the same session, loaded-artifact, diagnostic, and result objects. `promptabi plugins` makes those extension claims visible at the CLI before a user trusts them in CI. That prevents the common research-tool failure mode where each experiment has its own incompatible input format and output convention. The first workflow milestone also establishes practical CI affordances—discovery, overrides, caches, verbosity, deterministic rendering, explicit failure thresholds, and plugin-style check registration—before expensive formal checkers arrive.

The notebook surface makes that embedding path concrete for researchers and educators without adding a Jupyter dependency. `visualize_tokenization` returns the exact adapter token IDs, byte spans, added-token flags, special-token flags, and normalization steps. `visualize_template_rendering` pairs a concrete supported-fragment render with bounded symbolic paths and visible abstentions. `visualize_stop_reachability` shows stop strings and token IDs through the real tokenizer-backed stop analyzer. `visualize_grammar_product` summarizes a lazy DFA intersection search and shortest witness. `visualize_smt_constraints` reports the finite contract problem, backend, assignment or unsat core, checked-assignment count, and conclusion. `visualize_truncation` renders the existing budget visualization rows: kept or dropped segments, truncation boundary, overflow, and must-survive status. Each object implements `_repr_html_`, `_repr_mimebundle_`, `render_text`, `render_html`, and `to_dict`; the HTML is escaped and self-contained, so attacker-controlled prompt strings can be studied inside a notebook without becoming notebook HTML.

The corpus maintenance workflow now has the same integration discipline:

```
promptabi corpus manifest
promptabi corpus manifest --output seed-corpus.manifest.json
promptabi maintain refresh --output-dir maintainer_artifact --baseline previous_release
```

These commands validate the curated seed corpus without network calls, emit stable manifests with per-entry metadata hashes, tokenizer-config hashes, fixture hashes, expected behaviors, license notes, upstream references, pinned fixture revisions, and reproducibility notes, and fail deterministically if a required family or sentinel contract is missing. The maintainer refresh command extends that idea to the full release surface: it refreshes model-artifact manifests, structured-schema and provider-fixture manifests, the real-bug benchmark manifest, portable expected diagnostics with repository-relative spans, corpus diffs against a previous snapshot, and concise release notes. That turns the corpus from a collection of files into an auditable experimental input and release artifact.

12 Role-Boundary Non-Forgeability

The first major formal checker now starts from an executable role-boundary model and analyzes whether user, tool, function, or other attacker-controlled fields can render as structural prompt boundaries. In ChatML-like templates the relevant strings may be `<|im_start|>`, `<|im_end|>`, role names, or assistant prefixes. In Llama-style templates they may be header IDs and end-of-turn sentinels. In instruction-tag templates they may be `[INST]`, `[/INST]`, or custom fine-tune delimiters. In XML-like tool formats they may be tool-call start and end tags.

PromptABI already models the rendered prompt as regions: system, user, assistant, tool, developer, function, provider-specific symbolic roles, and generated assistant prefixes. The checker builds one global marker catalog for a chat template, split around placeholders so it never fabricates impossible concatenated markers. It then pairs every user-controlled content expression with every structural marker and separately checks raw dynamic role fields. Diagnostics identify the exact symbolic path and region, the controlled expression, the forged marker kind, a rendered excerpt showing where the forged boundary appears, and byte-level token evidence. The current implementation also records recognized sanitizers on role regions so JSON-encoded, escaped, URL/base64-wrapped, or sanitizer-preserving bound fields do not produce false-positive role-forgery reports.

This property is structural, not semantic. PromptABI can prove that a template allows user content to create an apparent assistant turn. It cannot prove that a model will obey that turn. That distinction makes the result more useful, not less, because the finding remains valid across model weights and sampling parameters.

The repository now exposes this distinction as an executable example rather than only a warning in prose. `examples/role-boundary` points at a ChatML-style template that directly interpolates `message['role']` and `message['content']`; the verifier fails with witnesses for raw `assistant`, `<|im_start|>`, and `<|im_end|>` markers. The paired `examples/role-boundary/safe.promptabi.json` uses the same control-token shape but JSON-encodes dynamic role and content fields, and the same check passes with no role-boundary diagnostics. Both artifacts are sha256-pinned and covered by focused CLI tests, so the documentation claim is exercised against real code.

13 Stop and Grammar Reachability

Stop policies often look simple: if the generated text contains a string, cut there. In structured-output systems this simplicity is dangerous. A stop string can be unreachable because the tokenizer or grammar cannot produce it. It can be ambiguous because multiple token paths or decoded byte strings correspond to the same visible text. It can overreach because the string appears inside a valid JSON string, tool argument, markdown fence, or XML-like field before the structured object is complete.

PromptABI now implements four concrete layers for this family: stop-policy parsing, tokenizer-backed stop alignment, fixture-backed stop differential replay, and bounded stop-overreachability. A dedicated parser accepts the shapes developers actually check into repositories or capture from serving stacks: OpenAI-compatible `stop` request fields, Hugging Face `generation_config` objects with `stop_strings` and EOS token IDs, llama.cpp/Ollama `reverse_prompt/antiprompt/ignore_eos` options, vLLM `sampling_params` with `stop_token_ids`, LiteLLM `litellm_params`, and nested framework kwargs such as `model_kwargs`, `generation_kwargs`, `parameters`, and `extra_body`. The loader normalizes those forms into a typed stop-policy artifact containing sorted unique stop sequences, stop token IDs, EOS behavior, source-family metadata, ignored empty stops, and JSON source spans. Malformed fields become deterministic loader errors rather than silent fallbacks.

The new tokenizer-backed layer analyzes every stop string under the selected tokenizer adapter and emits concrete byte/string/token alignment witnesses. It records UTF-8 bytes, token IDs, token surfaces, byte spans, normalized text, decoded text, round-trip metadata, multi-token segmentation, and special or added-token participation. It reports proper prefix and suffix collisions at the string, byte, and token-ID levels, normalization collisions where distinct configured stops collapse to the same tokenizer surface, and conservative unreachability findings for configured stop-token IDs that the chosen tokenizer cannot decode. The checker is intentionally careful not to claim that a string stop is unreachable merely because it does not round-trip through an encoder: many serving stacks match stop strings against decoded text streams. Those cases are reported as tokenizer-sensitive ambiguity instead, with heuristic modes.

The differential layer checks those static interpretations against recorded runtime traces without calling live providers. A provider-config fixture can store one or more `stop_trace` entries: generated text chunks, optional token IDs, the provider or framework family, and the expected stopped flag, matched stop, output text, finish reason, and whether the runtime retains the delimiter in the emitted output. PromptABI replays each trace through a small CPU-only simulator. OpenAI-compatible, vLLM, llama.cpp, Ollama, and LiteLLM-style families match raw substring stops on the cumulative decoded stream and exclude the stop delimiter from returned text; the Hugging Face stopping-criteria model records the common generated-text behavior where the stop string remains in the sequence that triggered stopping. Mismatches become deterministic diagnostics with replay witnesses, while agreement cases form an executable benchline. The repository ships both focused unit tests and a `examples/stop-policies` fixture that proves a vLLM-style trace agrees with the configured stop policy.

The overreachability layer separates structural delimiter collisions from supported content witnesses. Structural checks cover nested JSON delimiters, markdown code fences, XML-like tool-call tags, and OpenAI/Anthropic-style tool-call envelopes. Content checks are emitted only for supported JSON Schema and tool-parameter fragments where an unconstrained string field makes the stop text valid data, including bounded object and array-of-object parameter paths that occur in real tool schemas. All offsets are computed on the serialized wire text, including escaped characters, so the witness is not merely "the stop appears." It is a valid output prefix, the exact line and column where the stop fires, the parser state at truncation, and the malformed or prematurely accepted structure received by the application parser. PromptABI now analyzes the adjacent tokenizer x grammar emptiness product for bounded JSON Schema grammars, and should next combine stop policies with full grammars. A useful witness for grammar-level stop unreachability includes the stop sequence, its byte form, the tokenizer segmentation, and the grammar state that prevents production.

This check family is attractive because it is both practical and formal. It catches flaky production bugs, and it can be stated as a reachability property over a product of tokenizer, grammar, stop automaton, and parser states.

14 Structured Output and Grammar Emptiness

Structured-output libraries increasingly compile JSON Schema, regex fragments, EBNF-like grammars, Pydantic models, or provider-specific tool schemas into a decoder constraint. The application then parses the resulting bytes with a different parser. Bugs appear when those languages do not match. A schema may be valid under a Python validator but compile into an empty grammar under a specific backend. A grammar may accept strings the application parser rejects. A tokenizer may expose Unicode, escaping, whitespace, or added-token behavior that creates ambiguous structured outputs.

Tool schemas are now first-class inputs rather than opaque side files. PromptABI normalizes OpenAI `{type:function,function:{...}}` bundles, Anthropic `input_schema` tools, LangChain `args_schema` tools, Pydantic model JSON Schemas, TypeScript-style function arrays, MCP `inputSchema` tools, and provider response envelopes containing `tool_calls`. The result records the provider family, tool names, JSON-object versus JSON-string argument encodings, closed parameter schemas, required properties, enum and constraint paths, tool-call IDs, streaming chunk evidence, source spans, and ingestion issues such as open parameter objects. This gives later serialization, provider-migration, stop-overreach, and budget checks a typed schema substrate.

The next layer is now executable as `tool-serialization`. A recorded provider snapshot can declare the request tool names it accepts, the response tool-call names it emits, whether arguments are JSON objects or JSON-encoded strings, whether those strings are escaped, where tool-call IDs appear, whether parallel calls are enabled or observed, how streaming deltas fragment arguments, which parser names and encodings are accepted, and which delimiter belongs to the tool-call envelope. PromptABI compares that contract against the selected tool definition, chat template, stop policy, and parser declaration. It reports bounded diagnostics for name mismatches, argument-encoding mismatches, raw escaping risk, missing IDs, single-call parsers behind parallel providers, parsers that read each delta instead of assembling chunks, templates that do not render selected tools, and stop policies that omit the recorded tool-call delimiter. The fixture-backed test drives all of those findings through `promptabi verify`, with source spans into the provider JSON and witness steps for the compared fields.

The provider evidence is no longer an ad hoc collection of JSON snippets. PromptABI now ships a first-class provider fixture pack corpus. Each pack records the request surface, response envelope, tool-call name and argument paths, argument encoding, stop behavior, streaming delta assembly, error-code shape, and edge-case context/output limits for a representative provider family. The corpus loader validates those sections, materializes each pack as a loadable `provider-config` artifact, emits deterministic manifest hashes, and rejects credential-like keys or values such as authorization headers and API-key patterns. This gives provider-migration, tool-serialization, and future replay checks an offline input format that can be shared in bug reports and CI without storing secrets.

PromptABI now normalizes the supported JSON Schema subset into a typed tree that preserves objects, arrays, required fields, enum and const terminals, string and numeric constraints, type and composition unions, local reference expansion, bounded recursion abstentions, and additional-property semantics. It also compiles that tree into a bounded grammar IR with source maps, per-node parser states, a deterministic finite automaton over generated JSON witnesses, and round-trip checks against the real `jsonschema` Draft 2020-12 validator. The implemented `grammar-tokenizer-emptiness` check then pairs each local tokenizer artifact with each schema or grammar artifact. For the supported JSON Schema fragment it enumerates bounded DFA witnesses, encodes them with the real selected tokenizer adapter, decodes the resulting token IDs, and accepts only if the decoded text is still accepted by the grammar automaton. A satisfiable report records the JSON witness, token IDs, token surfaces, decoded text, grammar state path, and backend assumptions. An empty report shows the rejected candidate, such as a required "OK" literal rewritten by a lowercase-normalizing tokenizer into "ok". Under the stated bounded assumptions, if that product is empty, then the constrained-decoding backend cannot produce a valid object through that tokenizer path.

The implemented `grammar-tokenizer-ambiguity` check adds the adjacent question: even when a schema is non-empty, can the interface represent it in more than one structurally relevant way? PromptABI generates bounded schema-valid JSON values, expands them into lexical variants such as compact and spaced encodings, Unicode spellings, escaped strings, and added-token-friendly text, validates the variants against the real `jsonschema` backend when available, and then compares equivalence classes over token IDs, decoded text, and parsed structured values. It reports token-path conflicts when distinct schema values collapse to the same token sequence, decoded-text conflicts when normalization erases a distinction, byte aliases when multiple JSON byte spellings parse to the same value under different token paths, and added-token aliases when a single added token and its byte spelling decode to the same grammar text. Each finding carries both grammar texts, both token paths, both structured values, decoded text, and the bounded assumptions under which the ambiguity was proven.

Equally important is abstention. Arbitrary schemas with recursive references, custom validators, remote references, or backend-specific extensions should not be silently approximated as safe. The diagnostic should state exactly which fragment forced abstention and where that fragment appears.

15 Context-Window Budget Modeling

Token-budget failures are common because different layers count and truncate in different ways. A framework may drop old messages from the left. A RAG pipeline may trim chunks after adding metadata. A serving engine may reserve tokens for tools or generation prompts. A tokenizer mismatch may make a prompt fit in development but overflow in production. The result can be subtle: the user message survives, but the system instruction, output-format requirement, citation-bearing span, or tool schema disappears.

PromptABI now implements bounded must-survive verification over that modeling layer. A `prompt-segment` artifact names regions such as `system-policy`, `retrieval-context`, and `user-request`; marks regions as required or `must_survive`; carries declared `token_count` or `max_tokens` values; optionally includes content that can be counted through a real tokenizer adapter; and attaches per-segment overhead. The loader accepts explicit `{segments:[...]}` files and OpenAI-style message arrays, then merges file content with config-declared requirements so existing message fixtures become budget inputs without losing must-survive metadata.

A `framework-truncation-config` artifact contributes the maximum context window, output reservation, reserved tool tokens, generation-prompt tokens, model-specific special-token overhead, framework name, strategy label, model name, optional system/tool preservation flags, and role-priority drop hints. The `token-budget-model` check applies deterministic arithmetic:

```
input_budget =
    max_context
    - output_reserve
    - tool_reserve
    - generation_prompt_tokens
    - special_token_overhead
```

If a framework artifact and top-level config disagree about `max_context_tokens`, the artifact wins and the run emits an explicit conflict diagnostic. If required segments lack declared counts and no supported tokenizer can count their content, the checker abstains rather than inventing a tokenizer proxy. This preserves the repository's soundness boundary: arithmetic over declared or real-tokenizer counts is bounded and sound for the modeled inputs; estimates are not silently promoted into proof.

The check now normalizes framework-specific policies for LangChain, LlamaIndex, vLLM, Transformers, llama.cpp/Ollama, OpenAI-compatible servers, LiteLLM, common message-dropping strategies, and custom RAG pipelines. Left-truncating serving engines drop earliest prompt regions, conversation-memory frameworks drop oldest non-system messages by default, right and middle strategies are modeled explicitly, and RAG policies can prioritize retrieval-role loss before required instructions. Diagnostics report non-positive effective budgets, single segments larger than the input budget, required segment totals that cannot fit, total prompt overflow for optional content, optional segments dropped by the framework, and required segments that a concrete truncation policy can remove. For known counts and supported policies the checker also emits a **MustSurviveProof**: **proven** when every required region remains in the kept set, **violated** when truncation can remove one, and **abstained** when token counts or a custom policy fall outside the bounded model. Violations include a greedily minimized counterexample: the smallest ordered segment subset found by deletion that still makes the real policy drop a must-survive segment. The witness records the selected budget source, reserved-token arithmetic, known segment counts and their sources, unknown segments, kept segments, dropped segments, required-token sums, proof status, and counterexample token totals. The same diagnostic now carries a compact visualization row for every prompt segment: role, required/optional status, declared or tokenizer-derived count source, overhead contributions, cumulative token span, kept/dropped decision, truncation boundary, and survival label. This visualization is visible in human CLI text and is also emitted as the structured `token_budget_visualization` property in JSON and SARIF, so CI systems can annotate exactly which prompt fields disappear without scraping prose. The `examples/token-budget` workflow exercises this through the real CLI: a vLLM-style budget reserves output, tool, generation-prompt, and special-token overhead, then fails because required system and user segments exceed the modeled input budget and the selected truncation policy can remove required content.

The same prompt-segment path now carries RAG-specific chunk contracts rather than treating retrieved context as anonymous optional text. A retrieval chunk can declare its embedding or splitter tokenizer, document and chunk IDs, source character boundaries, packed prompt boundaries, expected and observed overlap, citation labels, citation-required flags, metadata token cost, prompt-template wrapper overhead, and a per-chunk retrieval payload limit. The `rag-chunking-compatibility` check reuses the budget simulator and emits bounded diagnostics when serving tokenizers disagree with chunking tokenizers, normalization or splitter drift changes chunk boundaries, citation-required chunks lose labels or are dropped by truncation, overlaps differ from the splitter contract, metadata dominates useful content, wrapper text pushes chunks beyond their configured size, or the retrieved payload cannot fit its declared limit. The `examples/rag-chunking` workflow drives those cases through the real CLI and records deterministic JSON witnesses, making RAG failures concrete rather than a vague "context quality" concern.

16 Differential Validation

PromptABI’s abstractions must be checked against the libraries developers actually use. The verifier should not implement a tokenizer, chat-template engine, or grammar compiler and assume equivalence. It should differentially test its abstraction against Hugging Face tokenizers, `apply_chat_template`, `tiktoken`, SentencePiece where practical, llama.cpp metadata, vLLM and OpenAI-compatible request fixtures, and structured-output backends such as Outlines, xgrammar, llguidance, lm-format-enforcer, Guidance, and Instructor.

The testing strategy has three tiers:

- unit and structured snapshot tests for public API stability, downstream plugin compatibility, deprecation metadata, generated API docs, and deterministic text, JSON, and SARIF rendering;
- tokenizer abstraction tests against real Hugging Face `tokenizers`, `tiktoken`, SentencePiece, and byte-level adapters;
- chat-template parser, symbolic-executor, concrete-renderer, role-region, and role-boundary non-forgeability tests across real `tokenizer_config.json` fixtures and known delimiter-collision regressions, including ChatML, Llama header tokens, Mistral instruction tags, XML-ish tool tags, markdown fences, fine-tune hash headers, generation-prompt extraction, dynamic role headers, content-control delimiter collisions, role assumptions, field references, special tokens, bounded message/tool loops, filters, whitespace controls, path budgets, empty tool lists, tool-role messages, and unsupported constructs, with concrete outputs differentially checked against Hugging Face’s renderer;
- loader and parser tests for local files, Hugging Face refs, tokenizer directories, GGUF stubs, provider snapshots, tool-definition schemas spanning OpenAI, Anthropic, LangChain, Pydantic, TypeScript-style, MCP, and provider tool-call envelopes, tool-serialization contracts over names, arguments, escaping, IDs, parallel calls, streaming chunks, template tool rendering, and stop delimiters, stop-policy configs spanning OpenAI-compatible, Hugging Face, llama.cpp/Ollama, vLLM, LiteLLM, wrapper kwargs, fixture archives, and source-mapped JSON fields;
- stop-tokenizer, stop-differential, and stop-overreachability tests that exercise byte/string/token alignments, multi-token stops, invalid stop-token IDs, normalization collisions, prefix/suffix collisions, added/special-token interactions, OpenAI-compatible and Hugging Face delimiter-retention semantics, provider fixture replay, CLI mismatch diagnostics, serialized stop offsets, parser-state truncation witnesses, supported schema/tool string witnesses, structural delimiter collisions, real-world Qwen/Gemma/llama.cpp stop bug reductions, and a real Hugging Face `tokenizers` BPE fixture;
- token-budget and RAG chunking tests that exercise context-window reservation arithmetic, required-segment overflow over declared counts, prompt-segment file/message merging, config-versus-framework budget precedence, LangChain/vLLM/custom-RAG truncation simulation, optional retrieval dropping, required segment truncation, order preservation, proven must-survive cases, minimized counterexample extraction, text/JSON/SARIF budget visualizations with token spans and dropped-field properties, real CLI diagnostics for output, tool, generation-prompt, and special-token overhead, plus tokenizer mismatch, chunk-boundary drift, citation loss, overlap accounting, metadata inflation, prompt-template overhead, and retrieval payload truncation over repository fixtures;
- tokenizer/config drift and configuration-diff tests that build baseline and current tokenizer directories, provider snapshots, and framework truncation artifacts, then verify special-token ID drift, added-token changes, normalizer changes, chat-template hash changes, BOS insertion changes, generation stop-policy deltas, provider-contract regressions, removed checks, context-budget regressions, and loader-abstention behavior through real CLI diagnostics;
- fixture and corpus tests for minimized tokenizer/chat-template bundles, public real-bug provenance, expected behaviors, sentinel consistency, artifact materialization, deterministic manifest generation, labeled structured-schema reductions, anonymized contracts, synthetic parser stress cases, tool-definition bundles, dynamic agent prompt-pack integration specs that generate real PromptABI configs, secret-free provider fixture packs with request/response/tool/stop/streaming/error/limit coverage and credential-pattern rejection, and real CLI replay of every included structured-schema config;
- future corpus tests for popular model and framework configurations pinned by immutable upstream revisions.

Differential tests are especially valuable for version drift. A tokenizer or provider SDK update can change behavior without changing application code. PromptABI should make that drift visible as an artifact contract change rather than a mysterious downstream parsing failure. The deployed stop differential harness is the first provider-fixture replay line: it is intentionally small, but it already distinguishes runtime families that remove a stop delimiter from those whose generated sequence still contains it.

17 Fixture Corpus

The fixture corpus is the bridge between research claims and developer trust. It should cover popular instruct-model templates, structured-output schemas, tool-call envelopes, provider fixtures, RAG budget cases, and known delimiter collision patterns. Each entry should include provenance, license notes, exact upstream revisions, hashes, minimized repro inputs, expected diagnostics, and unsupported-fragment notes.

The current repository now contains a curated tokenizer/chat-template seed corpus spanning Llama-style header tokens, Mistral instruction tags, Qwen and ChatML im-start/im-end variants, Gemma user/model turns, Phi role tokens, DeepSeek fullwidth control tokens, Zephyr role lines, OpenAI-compatible adapter envelopes, and Alpaca-like fine-tune delimiters. The entries are intentionally minimized: they retain the structural boundary tokens and template shape needed for static analysis while excluding model weights and heavyweight tokenizer vocabularies. Each entry records roles, sentinels, expected behaviors, license notes, fixture revision, upstream reference, reproducibility notes, and the fact that no download is required.

It also contains the first real-world bug corpus. Public reports from llama.cpp, vLLM, and Hugging Face Transformers are reduced to synthetic fixtures that preserve only the interface contract failure: Phi-style user content that can render role delimiters as structure, ChatML special tokens preserved by `apply_chat_template` as apparent control text, Qwen-style XML tool-call delimiters inside free-form tool parameters, streaming tool parsing interrupted by stop strings inside `function.arguments`, array-of-object tool parameters whose item strings can leak partial XML-style tool calls, multi-function XML boundary confusion, Gemma-style quote sentinels that decode to raw string delimiters, and a llama.cpp tool-call close marker that can stop before the downstream parser is in a safe state. Each case records its public GitHub reference, bug class, minimized artifact, and expected PromptABI witness. The tests load the corpus and assert that the executable role-boundary and stop-overreachability checkers produce those witnesses, making the paper claim auditable without network calls or copied upstream code.

The corpus is executable through the public API. `load_seed_corpus()` validates coverage of all required families, checks that declared sentinels appear in either the template or special-token metadata, verifies `add_generation_prompt` declarations against the template, computes sha256 hashes for metadata and tokenizer configs, and materializes entries as PromptABI tokenizer and parsed chat-template artifacts. The same path is exposed through `promptabi corpus manifest`, which produces a deterministic manifest for CI and paper artifacts.

The structured-schema layer is executable as well. It includes a reduced open-source-agent ticket-routing schema, an anonymized markdown-fence response contract, a synthetic XML tool-call parser-broader stress case, and an OpenAI-style tool-definition bundle that now runs through the real tool-schema-ingestion CLI check, and a synthetic tool-serialization contract that drives provider/tool/parser disagreement diagnostics through the real CLI. `load_structured_schema_corpus()` validates required source-category coverage, schema/grammar/tool entry types, labels, expected parser-compatibility status, per-entry artifact hashes, and runnable PromptABI configs. `promptabi corpus structured-schema-manifest` emits a stable manifest for release checks and paper tables. The next corpus layer should add exact immutable upstream revisions for vendored tokenizer configs, llama.cpp GGUF metadata, provider traces, and structured-output libraries.

The broader real-bug benchmark suite turns this fixture layer into an evaluation object. `fixtures/real_bug_benchmarks/benchmark` records seven labeled categories: popular template, tokenizer, tool schema, provider migration, structured-output library, RAG truncation, and training pipeline. The loader validates category coverage, provenance, labels, and expected rule IDs, then replays each entry through the same public analyzers used by the CLI: role-boundary non-forgeability, stop-overreachability, tokenizer differential comparison, provider migration diagnostics, structured-schema parser compatibility, RAG chunking verification, and finite static contracts with the solver-independent enumeration backend for training target-role and supervised span-region alignment, source leakage, and tokenizer/template stage drift. The suite therefore measures whether PromptABI still detects a diverse labeled bug set, not merely whether a fixture parses.

The corpus should avoid secrets and network requirements. Provider fixtures should be recorded request and response shapes, not live API calls. Model artifacts should be metadata and tokenizer files, not weights. This keeps CI fast, reproducible, and privacy-preserving.

18 Benchmark Lines

PromptABI now ships an executable CPU-only benchmark suite rather than a single smoke loop. The entrypoint `benchmarks/benchmark` delegates to a package-level harness that returns stable JSON records with benchmark name, iteration count, runtime, throughput, and case-specific metrics. Every case runs against real PromptABI code and checked-in examples or fixtures; no network, model weights, or provider credentials are required.

Benchmark	Measurement
Tokenizer analysis	Byte-level encode/decode metadata extraction across JSON-like payloads, Unicode normalization, added tokens, special tokens, and exact round-trip counts.
Template symbolic execution	Ten seed-corpus Hugging Face <code>tokenizer_config.json</code> templates parsed and boundedly symbolically executed with path and abstention counts.
Grammar emptiness	A real structured-schema fixture compiled to the bounded JSON Schema DFA product and checked against the byte-level tokenizer with state, accept, and candidate counts.
Stop checks	Stop-tokenizer alignment plus bounded stop-overreachability over structural JSON, markdown, XML-like tool-call, and provider-envelope regions.
Z3 static contracts	Declarative interface rules for allowed roles, required regions, forbidden delimiters, schema obligations, stop policies, linted impossibility/vacuity/fragment/suppression hazards, layered policy/app/pack/training composition, and finite budget invariants lowered to the existing SMT layer or finite fallback.
Budget checks	Segment accounting, tokenizer-derived content counts, vLLM-style reservations, truncation simulation, required-token overflow, and finding counts.
Corpus verification	<code>promptabi corpus verify</code> release-gates the seed corpus, structured-schema corpus, provider fixture packs, real-bug benchmark, eval fixture pack, labeled evaluation, minimized SMT benchmark, and optional runtime/Python-heap thresholds with explicit nonzero coverage checks.
Minimized SMT benchmark	<code>promptabi corpus smt-benchmark</code> replays satisfiable, unsatisfiable, timeout-prone, and unsupported reduced prompt-interface formulas against the real finite/Z3 solver and emits stable query hashes, checked-assignment counts, conclusions, and manifest hashes.
Compatibility audit	<code>promptabi release compatibility-audit</code> composes the real corpus gate, provider fixture replay, grammar differential replay, and parser/framework compatibility fixtures for tokenizer, template, provider, grammar, and framework candidate-version surfaces, returning pass/fail/abstained evidence instead of treating an unpinned version label as proof.
Maintainer refresh	<code>promptabi maintain refresh</code> regenerates model/corpus/provider manifests, expected diagnostic baselines with portable fingerprints, corpus diffs, and release notes from real repository configs.
Cold/warm cache	Direct measurement of token-budget analysis before and after per-session cache reuse, including an explicit cache-reused metric.
Real-bug benchmark	<code>promptabi corpus real-bug-benchmark</code> validates and replays the labeled real-bug suite across popular templates, tokenizers, tool schemas, provider migrations, structured-output parsers, RAG truncation, and training manifests, including target-role, source-leakage, synthetic-generator, and stage-drift contracts, emitting observed rule IDs and a stable manifest hash.
Public bug gallery	<code>promptabi corpus bug-gallery</code> reruns the same benchmark cases, then emits a public-facing text, JSON, or Markdown gallery with sanitized artifact digests, minimized repro metadata, replay hashes, root-cause explanations, concrete fixes, and public GitHub patch or issue links without copying upstream source or private payloads.
Beta report	<code>promptabi corpus beta-report</code> replays open-source-style beta case studies, upstream issue records, tuned-absent false-positive expectations, and a synthetic abstention-message regression, requiring actionable abstention witnesses rather than success-shaped unknowns.
Labeled evaluation	<code>promptabi corpus evaluation</code> runs the default evaluation corpus, expands the real-bug benchmark and eval fixture pack into per-case labels, measures scoped precision/recall/F1, abstention rate, runtime, peak memory, witness quality, Z3-backed solver decisiveness, and differential agreement while keeping unrelated informational diagnostics out of false-positive counts.

The focused benchmark tests execute selected cases against repository fixtures, assert that the registry covers every deployed surface, validate the JSON schema of the benchmark output, reject unknown benchmark names, and run the CLI wrapper. A one-iteration full sweep currently reports, among other values, ten templates and 163 symbolic paths for the template case, 58 grammar states for the structured-schema grammar product, five stop-overreachability findings for the stop case, two SMT findings for the static-contract case, 32 diagnostics across six corpus configs, and a true cache-reuse observation for the warm path. These numbers are not final performance claims; they are reproducible benchlines that will make future optimization, regression detection, and paper artifact evaluation concrete.

19 Documentation and Developer Experience

A 1000-star verification repo needs more than algorithms. It needs a first-run experience where a developer immediately understands what failed and how to fix it. The README now leads with a runnable failing ChatML diagnostic, install commands, GitHub Actions wiring, a CPU-only/non-inference boundary, and the guarantee taxonomy for role boundaries, stops, grammars, tools, providers, budgets, RAG, training, and evaluation contracts. The docs tree explains architecture, quickstart, check families, the artifact model, and a focused role-boundary page that separates structural non-forgability from semantic prompt-injection resistance while linking to the runnable unsafe/sanitized example. It also ships a security model that tells users exactly what is local, what is sensitive, what is outside scope, how provider fixtures and solver inputs are handled, and how to prepare a responsible upstream disclosure. The contribution guide states the central norms: deterministic output, CPU-only tests, minimized fixtures, typed public APIs, and explicit soundness boundaries.

The examples directory is equally important. It now contains the minimal typed-artifact workflow, role-delimiter injection, token-budget overflow, RAG chunking, provider migration, structured-output, and a chat SFT quickstart that checks rendered roles, token spans, packing, loss masks, stage pins, and redacted witnesses without model weights, plus dynamic agent-framework examples that materialize prompt-pack runtime assembly into verifiable configs and catch tool, stop, and provider drift before deployment. PromptABI should keep adding small buggy/fixed applications for stop overreach, provider AST drift, tokenizer drift, and training/eval contracts; each pair gives users a recipe and maintainers a stable regression.

The CLI remains terse in CI and expandable locally: `promptabi explain` adds witnesses and fixes, `promptabi bug-report` emits sanitized upstream markdown, and `-local-summary` records only aggregate local usage.

20 Evaluation Design

The evaluation answers five questions.

1. Does PromptABI find real interface bugs in popular LLM application stacks?
2. Are the findings precise enough that developers would act on them?
3. Does the verifier abstain honestly when artifacts exceed the supported fragment?
4. Do automata and SMT-backed finite contracts produce reproducible witnesses?
5. Is the runtime low enough for pull-request CI?

The deployed setup includes `fixtures/evaluation/labeled_corpus.json`, `fixtures/evaluation_fixture_packs/pack.json`, and `promptabi corpus evaluation`, which execute real configs plus every replayable real-bug and eval-fixture case and report precision/recall/F1 over an explicit per-case rule universe. The training slice includes RLHF/DPO fixtures, the SFT quickstart, and a multi-adapter data-loader fixture pack covering Hugging Face Datasets, JSONL, Parquet metadata, Axolotl, LLaMA-Factory, TRL, OpenRLHF, internal redacted schemas, and bounded streaming JSONL shards; clean manifests prove safety, while defective manifests trigger static counterexamples for source leakage, prompt-hash, mask-policy, truncation, role-layout, response-start, packing-boundary, and synthetic generator role/schema/tool/truncation bugs. The evaluation slice now also checks HELM/lm-eval/custom/provider-hosted tokenizer pins for drift or abstention, replays fixture-pack cases for stop strings, parsers, truncation, role-boundary forgery, and tokenizer mismatch, proves that grading-parser assumptions agree with constrained decoding schemas and application output parsers over bounded witnesses, and proves that answer-key, private rubric, and chain-of-thought variables do not overlap model-visible prompt variables, with `examples/evaluation-harness/unsafe.promptabi` exercising all three leakage diagnostics against real source-mapped JSON. The same report measures abstentions, runtime, memory, witness completeness, Z3-backed solver decisiveness, and differential agreement against recorded real-library semantics. Published eval scores also need a replayable contract freeze rather than only a metrics table. PromptABI now includes `promptabi corpus evaluation-reproducibility`, a harness-driven mode that loads real `evaluation-harness` configs and emits stable hashes for exactly the surfaces that determine a benchmark run: prompt-rendering fields plus the chat-template artifact hash, benchmark-tokenizer snapshots plus serving tokenizer pins, provider fixture metadata and hashes, stop-policy sequences and token IDs, and parser/schema contracts. The manifest fails incomplete when any of those five surface families is missing, and a real stop-policy mutation changes the surface hash in tests. This gives paper tables, release notes, and downstream eval dashboards a compact artifact that says not only “these metrics passed”, but “these prompt, tokenizer, provider, stop, and parser contracts were the ones actually pinned.” The companion `promptabi corpus beta-report` command answers a different question: whether beta evidence is maintainable. It replays open-source-style case studies, keeps upstream issue records beside the configs that reproduce them, checks tuned-absent rule IDs so false-positive regressions are visible, and requires abstentions to include actionable witnesses and setup suggestions. The maintainer-facing `promptabi corpus verify` command composes this evaluation with seed-corpus manifests, structured-schema config replays, provider fixture replay coverage, real-bug and eval-fixture replay, and optional runtime/Python heap thresholds. It fails closed on missing coverage, disappearing expected rule IDs, low witness quality, missing differential or Z3-backed coverage, or exceeded performance ceilings, making release-blocking soundness and performance regressions visible through one JSON or text report. Future experiments can add mutation-based stress cases and version-drift tests. The deployed seed corpus already supplies the first corpus-wide experimental unit: a run can validate ten families, materialize twenty PromptABI artifacts, and compare the resulting manifest hash across revisions. Future experiments can swap minimized fixtures for revision-pinned upstream tokenizer configs while keeping the same manifest schema.

The structured-schema corpus supplies the first labeled schema/tool-call compatibility unit: tests replay expected agreement or mismatch labels through real schema and regex artifacts, execute every fixture’s PromptABI config through the CLI, and produce a deterministic manifest containing source categories, expected rule IDs, hashes, and provenance.

The real-world bug corpus supplies the first labeled detection experiment. In the current snapshot, eight public bug patterns map to executable checks: Phi-style role delimiter forgery from llama.cpp, special-token injection through Hugging Face `apply_chat_template`, Qwen XML tool-call delimiter overreach from vLLM, streaming stop interruption inside OpenAI-compatible `function.arguments`, llama.cpp array-of-object tool-argument leakage, vLLM multi-function XML boundary confusion, Gemma quote-sentinel truncation from vLLM, and a llama.cpp tool-call parser-boundary stop failure. PromptABI detects all eight with bounded structural witnesses. The result is intentionally modest but important: the evaluation is no longer only synthetic, and each claimed detection is backed by a minimized fixture that can run in ordinary CPU-only CI.

The new real-bug benchmark suite promotes those reductions into a cross-surface benchmark line. It requires all seven failure categories listed above, replays every label against real code, records observed rule IDs, and fails manifest generation if any expected rule disappears. This creates a release-blocking regression target for practical prompt-interface failures: a tokenizer oracle drift, a provider migration mismatch, a parser broader-than-schema bug, a RAG chunking/citation loss, and a discrete training target-role violation are tested beside public template and tool-stop reductions.

The public bug gallery is the human-facing counterpart to that manifest. It is not a hand-edited marketing page: `promptabi corpus bug-gallery` loads the same benchmark suite, replays each case, hashes the sanitized replay payload, keeps only minimized structural witness metadata, and attaches root-cause text, fixed-version/workaround guidance, and public GitHub issue or patch links. The `launch-assets` command consumes that exact report, so the website-facing gallery and the release-blocking benchmark cannot drift apart silently.

The grammar-emptiness benchline adds a complementary satisfiability experiment: one fixture proves that an ordinary byte-level tokenizer can produce the bounded JSON Schema object witness {}, while another proves that tokenizer normalization can make a const-schema product empty. Both cases run through the same public API and CLI diagnostic contract rather than private test helpers.

The strongest paper result would be upstreamable bugs: minimized reports that maintainers of templates, adapters, structured-output libraries, or framework integrations accept as real. That evidence is more compelling than a large synthetic benchmark alone. PromptABI now has the local machinery for that evidence: a selected diagnostic can be rendered as a markdown issue with the affected artifact table, pinned hashes or revisions, reproduction commands, expected and actual behavior, redacted witness steps, minimized repro metadata, fix suggestions, and an explicit no-network privacy note. The current tests generate such an issue from the real unsafe ChatML role-boundary fixture, so upstream reporting is an exercised workflow rather than a paper-only plan.

21 Threat Model, Limitations, and Non-Goals

PromptABI’s threat model is structural. Attacker-controlled content may enter user messages, tool outputs, retrieval chunks, metadata, or provider-compatible fields. The verifier asks whether that content can be represented as control structure, whether output can be truncated in a way that changes parser semantics, whether a requested structured-output language is empty or ambiguous, and whether required prompt segments survive budget policies.

The implemented security model makes the trust boundary explicit. PromptABI runs on a local developer machine or CI runner, treats prompt-interface artifacts and rendered witnesses as sensitive, and does not require live providers, model weights, GPUs, or remote solvers for the core verification claims. Provider fixture packs are required to be deterministic, anonymized or explicitly labeled, download-free, and secret-free; the loader rejects credential-like keys and values such as API keys, authorization headers, bearer tokens, cloud access keys, and passwords before those fixtures become benchmark evidence. Solver-backed checks lower bounded strings, enums, lengths, parser states, role regions, and budget facts into local Z3 formulas or finite enumeration; those formulas are treated as local artifacts, not telemetry. Usage summaries are opt-in local JSONL records containing command names, durations, exit codes, and aggregate counts, and the CLI privacy command now states that prompts, schemas, configs, constraints, witnesses, artifact contents, file paths, and network sends are excluded.

Artifact privacy also shapes the reporting workflows. SARIF, HTML, text, JSON, and markdown bug reports are local outputs that may contain source spans or structural witnesses and should be handled like build artifacts. The `promptabi bug-report` path generates sanitized, length-limited markdown with pinned artifact metadata and an explicit no-network privacy note, but it does not file issues or transmit data. The recommended responsible-disclosure path is therefore reproduce locally, explain the diagnostic, minimize the artifact where possible, review the sanitized report for private names or business logic, and coordinate privately first for vulnerabilities that could expose credentials, cross-tenant data, unsafe tool execution, or broadly deployed prompt-interface bypasses.

The non-goals are equally important. PromptABI does not prove that a model will resist prompt injection. It does not prove that a model will choose a safe tool. It does not evaluate factuality, alignment, toxicity, or behavioral compliance. It does not require model weights or generation. If an artifact contains arbitrary code, unsupported Jinja, provider behavior not captured by fixtures, or schemas outside the supported fragment, the correct behavior is to abstain or emit a bounded finding with explicit assumptions.

This narrowness makes the tool easier to trust. A structural counterexample does not depend on random seeds, temperature, model provider, or GPU hardware. If the artifact composition admits the unsafe state, that state exists.

22 Related Systems

PromptABI is adjacent to but distinct from model execution verifiers, prompt-injection scanners, schema validators, grammar-constrained decoding libraries, and provider SDK test suites. Model execution verifiers focus on tensors, devices, dtypes, shapes, checkpoint compatibility, or export gates. Prompt-injection scanners often attempt behavioral or semantic classification. Schema validators check a schema against an instance, but not necessarily the product of schema, tokenizer, grammar backend, stop policy, and application parser. Constrained-decoding libraries enforce output languages but may not report whether their language matches the downstream parser or whether stop logic truncates inside that language.

PromptABI’s niche is the composition boundary. It verifies that tokenizer, template, tool, grammar, provider, and budget artifacts agree. It should integrate with existing libraries rather than replace them: validators provide ground truth for schema fragments, tokenizer libraries provide differential oracles, and provider fixtures define offline request/response semantics.

The closest analogy is a static analyzer for ABI compatibility. The concern is not whether either side of an interface is individually well-written; it is whether their composed contract remains valid as artifacts evolve.

23 Roadmap Without Step Enumeration

The roadmap is best understood as a progression from stable surfaces to formal checks to ecosystem integration. The skeleton, core artifact model, loader layer, diagnostic contract, CLI, tokenizer abstraction, finite automata layer, finite-state transducer layer, finite-contract solver, source-span tracking, differential tokenizer and stop-policy harnesses, the curated seed corpus, the manifest pipeline, and the first sanitizer-aware role-boundary checker have begun the experimental layer. The next phase should deepen stop-reachability, grammar-emptiness, and budget-survival checkers. Corpus growth, lockfiles, drift detection, SARIF, GitHub Actions, pre-commit hooks, HTML reports, editor `publishDiagnostics` payloads, suppression policies, and plugin APIs turn the verifier from a library into an everyday CI and IDE tool.

The research roadmap adds executable specifications, proof sketches for supported fragments, richer formal-methods documentation, and post-release corpus audits on top of the shipped reproducibility packages, mutation fuzzing, cross-version CI, release automation, and real-bug benchmark suites. The community roadmap adds contribution infrastructure, compatibility matrices, integration guides, launch assets, and governance. These activities are not separate from the technical work. They are how a formal verifier becomes durable enough for production teams to trust and for researchers to reproduce.

24 Current Smoke Results

The current codebase is validated by focused executable tests rather than a single end-to-end smoke. They cover the typed package, artifact loaders, source-span mapping, diagnostics, renderers, public API, CLI, tokenizer adapters, formal automata/transducer/SM primitives, corpus manifests, provider fixtures, budget/RAG checks, dependency graph visualization, release gates, and paper artifact. The table below names the concrete benchlines that matter for reproducibility; each row is backed by checked-in fixtures or examples and runs without GPUs or provider calls.

Line	Current status
Focused tests	Artifact model, public API, config loading, CLI behavior, dependency graph rendering, source spans, diagnostics, and render snapshots pass in isolation.
CI/release tests	Workflow fixtures cover Python/OS/backend matrices, optional dependency slices, semantic-version gates, PyPI builds, standalone executables, Docker images, signed release assets, and changelog checks.
Formal core tests	Automata reachability/difference, transducer composition/projection, over-approximation, finite-domain solving, unsat cores, and Z3-backed assignment extraction pass against executable code.
Check-mode taxonomy tests	Diagnostic objects, JSON output, text output, and SARIF output preserve sound/complete/bounded/Z3-backed/heuristic/abstaining metadata deterministically.
Tokenizer/template tests	Byte-level normalization, added/special tokens, Hugging Face ByteLevel BPE, <code>tiktoken</code> , SentencePiece, and ten-family Hugging Face chat-template symbolic execution match real backend behavior.
Role-boundary checker tests	Bounded non-forgeability catches dynamic role-header and delimiter forgery across ChatML, Llama, Mistral, XML tool tags, markdown fences, and fine-tune hash headers with minimized witnesses and sanitizer-aware false-positive suppression.
Real-world bugs	Eight public llama.cpp/vLLM/Hugging Face bug patterns plus the real-bug benchmark suite are reduced to synthetic offline fixtures, detected with concrete witnesses, and rendered as a sanitized public bug gallery with replay hashes, fixes, and upstream patch links.
Stop/provider/structured-output tests	Recorded provider chunks replay through CPU-only simulators; parser compatibility, tool-schema ingestion, tool serialization, grammar emptiness, grammar ambiguity, and stop overreachability run through real CLI fixtures.
Evaluation-harness tests	<code>examples/evaluation-harness/*.promptabi.json</code> prove provider/model/tokenizer/template/grading-parser/application-parser/stop/few-shot compatibility, catch eval-tokenizer drift, and reject private eval fields rendered into prompts.
Corpus commands	<code>promptabi corpus manifest</code> , <code>promptabi corpus structured-schema-manifest</code> , <code>promptabi corpus real-bug-benchmark</code> , <code>promptabi corpus bug-gallery</code> , <code>promptabi corpus evaluation-fixture-pack</code> , <code>promptabi corpus smt-benchmark</code> , <code>promptabi corpus verify</code> , <code>promptabi corpus evaluation</code> , and <code>promptabi corpus beta-report</code> validate fixture coverage, metrics, witness quality, hashes, and CLI/public-API renderers.
Budget/RAG/drift tests	<code>examples/token-budget/promptabi.json</code> and <code>examples/rag-chunking/promptabi.json</code> cover must-survive segments, framework reservations, truncation decisions, citation loss, overlap errors, metadata inflation, tokenizer mismatch, and config/tokenizer drift.
Benchmark and maintainer lines	<code>benchmarks/benchmark_smoke.py</code> emits JSON for cold/warm cache and checker families; <code>promptabi maintain refresh</code> writes seed, structured-schema, provider-fixture, real-bug, corpus-verification, expected-diagnostic, diff, snapshot, and release-note artifacts.
Verified configuration gallery	<code>promptabi gallery</code> reruns five curated configurations covering pinned OpenAI-style tools, sanitized ChatML role boundaries, training-to-serving static contracts, RAG truncation budgets, and audited accepted-risk suppressions; focused tests assert deterministic JSON/text renderers, actual solver-backend summaries, stage-pin evidence, lockfile/pin badges, and risk explanations.
Supply chain, editor, API, docs	<code>artifact-provenance</code> , <code>promptabi diagnostics lsp</code> , <code>promptabi api-docs</code> , MkDocs navigation, and security-model tests cover provenance, LSP ranges, public API compatibility, privacy, <code>beta-report</code> , minimization, and provider-fixture safety

These are not the final experiments. They are the baseline benchlines that later experiments will extend: every new checker should add focused tests, fixture cases, corpus cases where relevant, and a benchmark line that can be compared across revisions.

25 Why CPU-Only Is a Strength

Many LLM quality questions are inherently behavioral and require model access. PromptABI intentionally avoids those claims. Tokenization runs on CPU. Chat-template rendering runs on CPU. JSON Schema normalization, grammar compilation, automata and transducer operations, parser checks, source-span tracking, fixture replay, and truncation simulation run on CPU. Provider compatibility can be tested from recorded fixtures. Context-window analysis depends on token counts and framework policies, not logits.

This makes PromptABI deployable in environments where GPU inference, external API calls, and prompt upload are not acceptable. A company can run the verifier inside CI on private prompt artifacts. A maintainer can run corpus checks on a laptop. A paper artifact can be reproduced without model weights. A pull request can fail because a tokenizer config changed, not because a stochastic generation happened to expose the bug once.

The CPU-only boundary also sharpens the formal claims. The tool can say "this delimiter is forgeable" or "this stop sequence can truncate valid JSON" with a structural witness. It should not say "the model is safe." That honesty is a competitive advantage.

26 Conclusion

PromptABI verifies a fragile and under-tooled layer of LLM systems: the discrete interface between application data structures, prompt rendering, tokenization, constrained output, provider protocols, truncation policies, and application parsing. The repository now combines a typed package, formal automata and SMT core, deterministic CLI and API surfaces, source-mapped diagnostics, minimized witnesses, ranked fix triage, real fixture corpora, precision/recall evaluation, CPU-only benchmarks, mutation-fuzzing experiments, a verified configuration gallery, policy gates, GitHub integration, and a generated paper artifact with frozen hashes and regeneration commands.

The research foundation is the transducer view of prompts, tokenizers, grammars, stop policies, parsers, and budgets. If the project continues to preserve deterministic outputs, honest abstention boundaries, replayable witnesses, and real corpus validation, it can become both a practical CI tool and a credible systems research contribution. The core insight is simple but powerful: many production LLM failures happen before the neural network is involved, and those failures can be found before deployment.